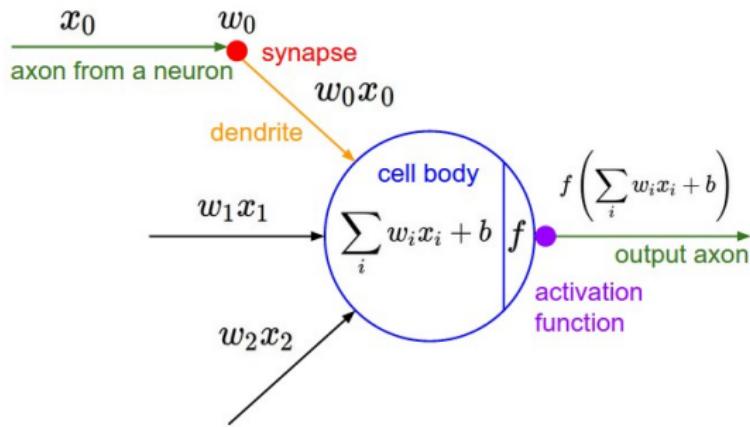


Primer on ML in HEP [3]

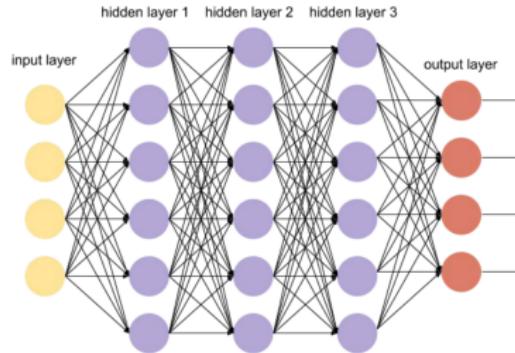
Overview

- Jets
 - b-tagging
 - t-tagging
 - grooming&PU
- Going adversarial
 - mass sculpting
 - DeepSF
 - domain adaptation
- Simulating with GAN&VAE
- Trigger
- Tracks
- * Bonus
 - Kaggle
 - INFERNO

NN [recap]



NN [recap]

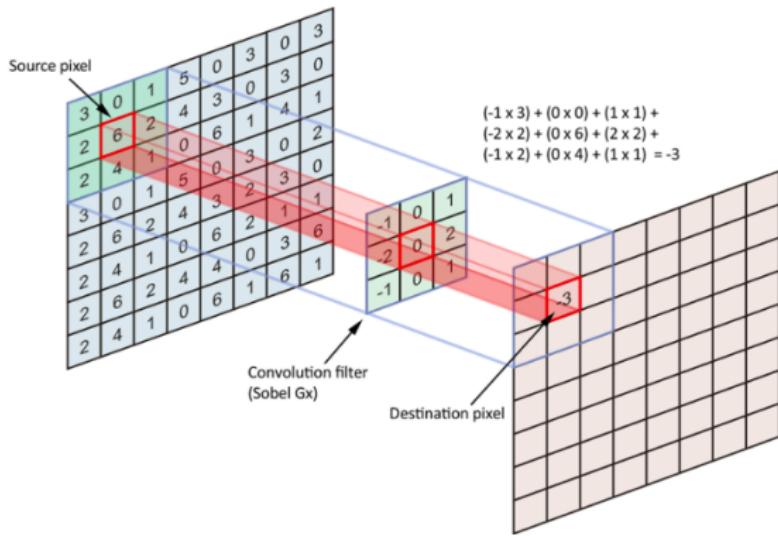


$$h_t = f(Wh_{t-1} + b)$$

$$a(x) = f_3(W_3f_2(W_2f_1(W_1f_0(W_0x + b_0) + b_1) + b_2) + b_3)$$

CNN [recap]

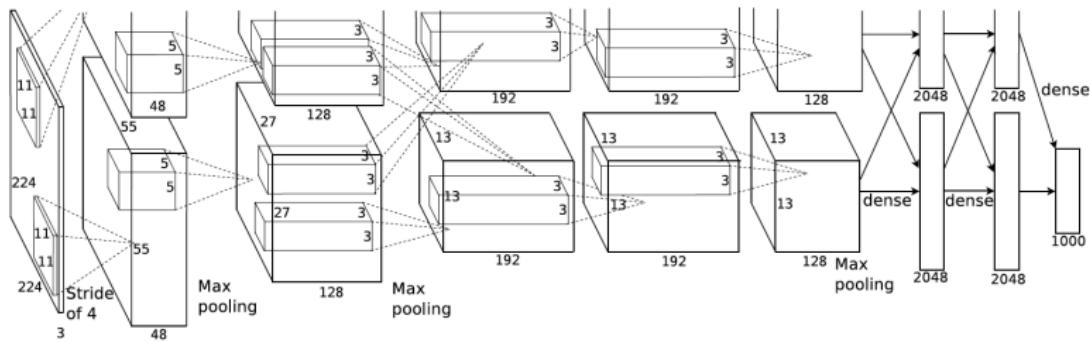
conv. layer



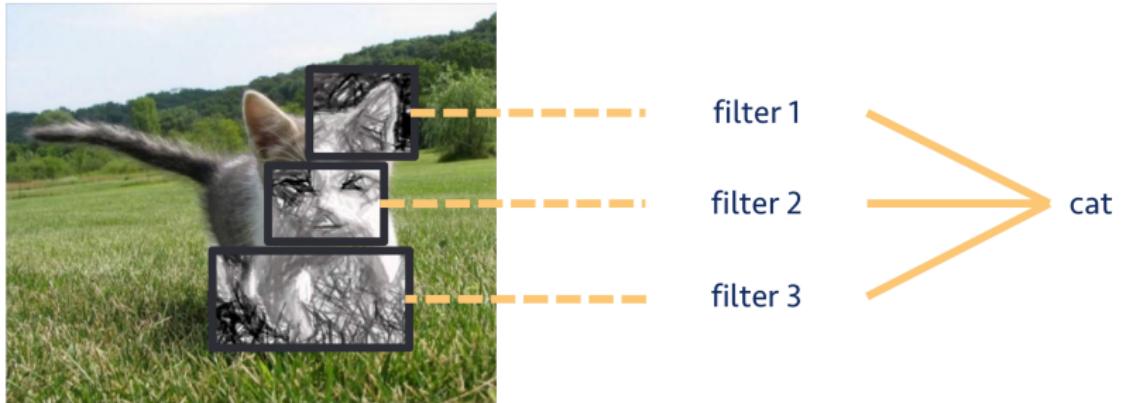
from [paper]

CNN [recap]

AlexNet

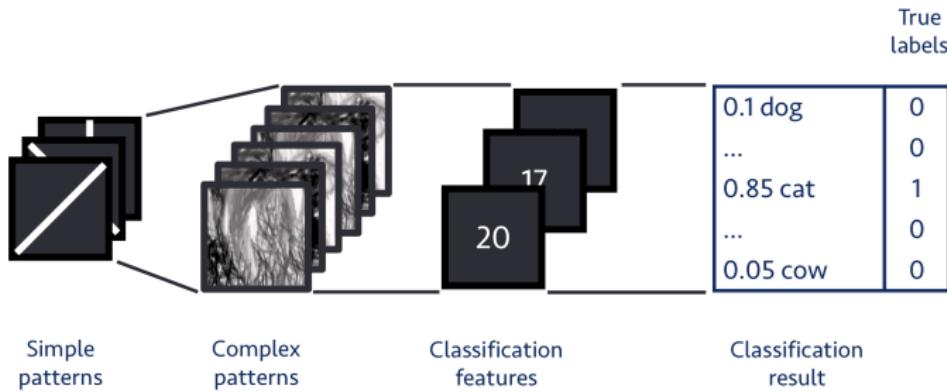


CNN [recap] patterns



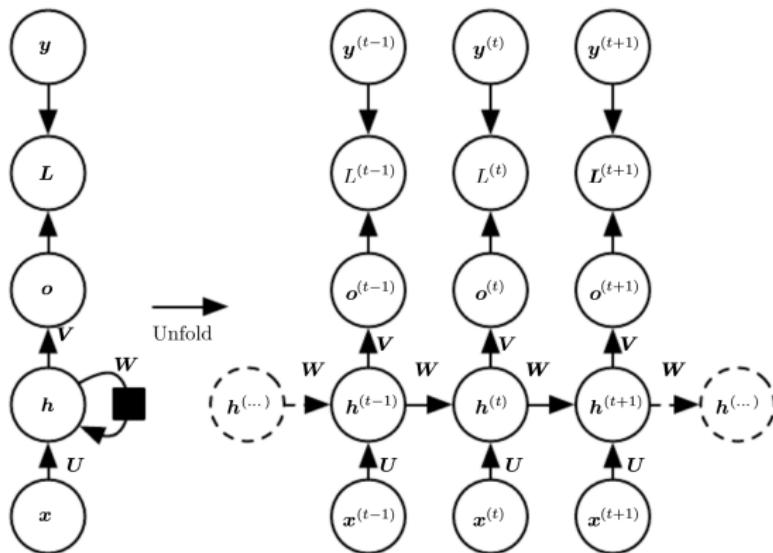
CNN [recap]

patterns

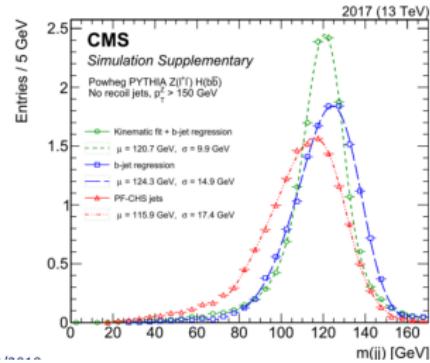
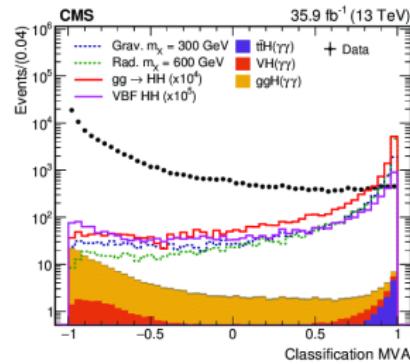


RNN [recap]

vanilla edition



- classification
 - μ, e^-, τ ID
 - sig vs. bkgr
- regression
 - mass resolution ($H \rightarrow bb$)



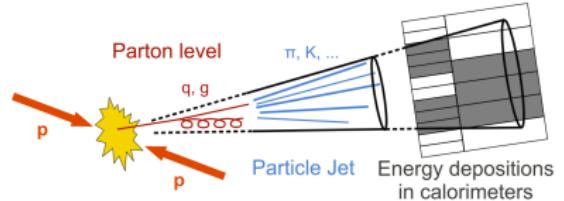
Jets

tagging

Tagging

why bother

- distinguish btw. b/c/uds/g
- tag t/W/Z/H
- measure SM (e.g. VH(bb), ttH, H → cc)
- search for new (X → bb)

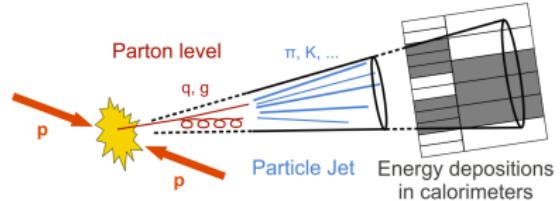


Tagging

why bother

- distinguish btw. b/c/uds/g
- tag t/W/Z/H
- measure SM (e.g. VH(bb), ttH, H → cc)
- search for new (X → bb)

How would you represent a jet?

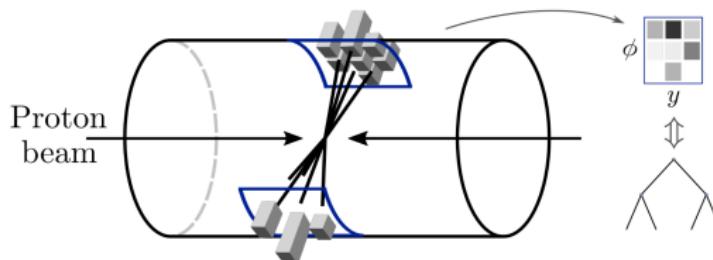


What's a jet, anyway?

A jet can be described in a variety of ways

- ▶ Sum of its particle constituents
- ▶ Image of energy deposits in rapidity-azimuth
- ▶ Complete basis of observables (e.g. energy flow polynomials)
- ▶ Image of primary Lund declustering sequence
- ▶ Binary tree associated with the jet clustering sequence
- ▶ ...

Most of them involve some trade-off between interpretability, representability, robustness and performance.

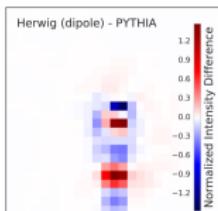
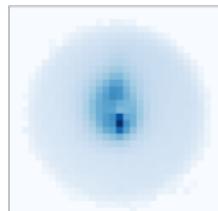
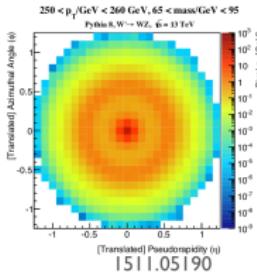


Frédéric Dreyer

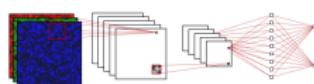
4/14

[1]

AS IMAGES...

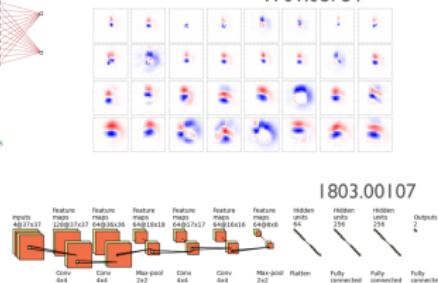


ATL-PHYS-PUB-2017-017

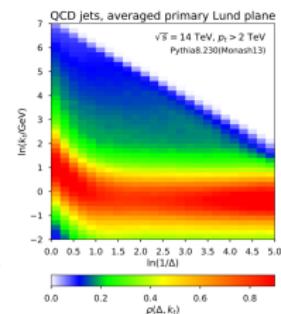


red = transverse momenta of charged particles
green = the transverse momenta of neutral particles
blue = charged particle multiplicity

1612.01551



1803.00107



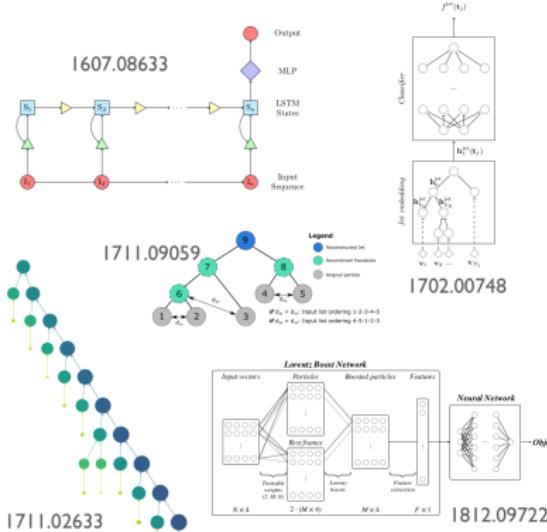
1807.04758

And more...

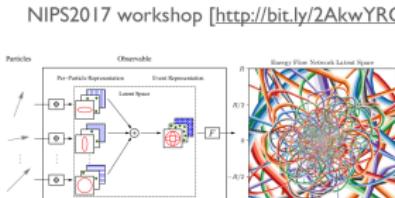
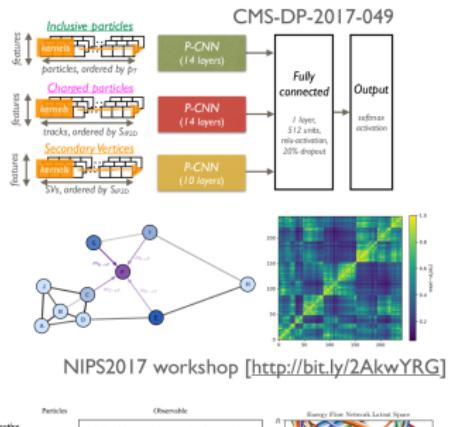
3

[2]

As COLLECTIONS OF PARTICLES...



$$\text{with } C = \begin{pmatrix} 1 & 1 & \cdots & 0 & \chi_1 & \cdots & 0 & C_{1,N+2} & \cdots & C_{1,M} \\ 1 & 0 & \cdots & 1 & 0 & \cdots & \chi_N & C_{N,N+2} & \cdots & C_{N,M} \end{pmatrix}, \quad \bar{k}_j \xrightarrow{\text{LoD}} \bar{k}_j = \begin{pmatrix} m^2(\bar{k}_j) \\ p_T(\bar{k}_j)/\Delta R_{j,\text{jet}} \\ w^{(j)} E(\bar{k}_{j\text{m}}) \\ w^{(j)}_{\text{min}} w^{(j)}_{\text{max}} \\ E_T(\bar{k}_j) E_T(k_{\text{m}})/(\Delta R_{j\text{m}})^{0.2} \end{pmatrix}.$$



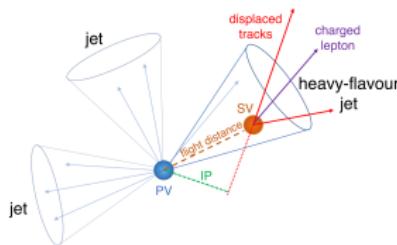
1810.05165

And more...

4

[2]

Tagging CMS taggers



circa 2017

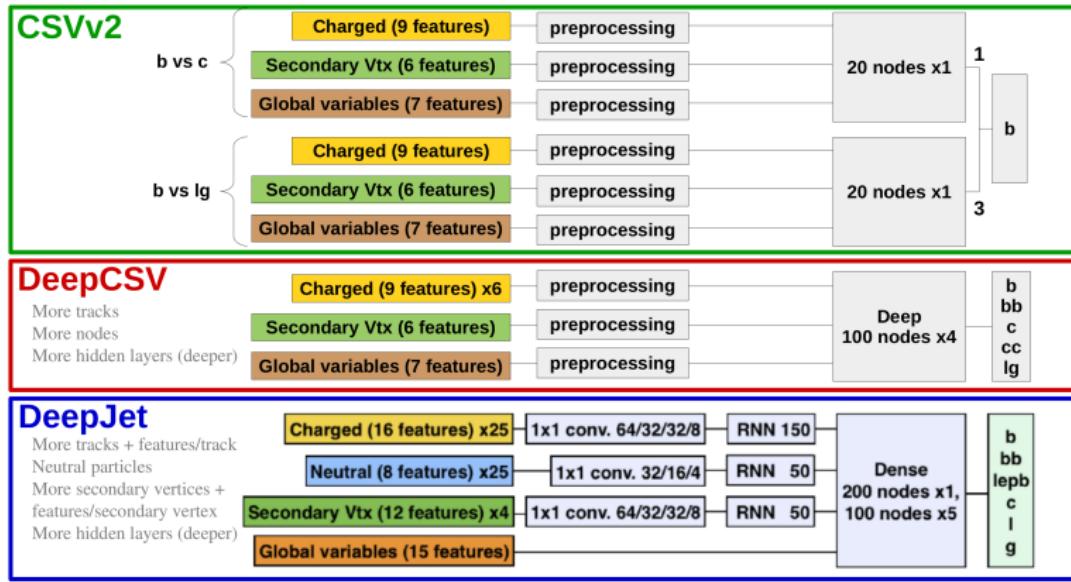
- JP/JBP
- CSVv1/2
- SE/SM
- cMVAv2

circa 2019

- DeepCSV
- DeepJet
- DeepDoubleX
- DeepAK8
- DeepYouNamelt

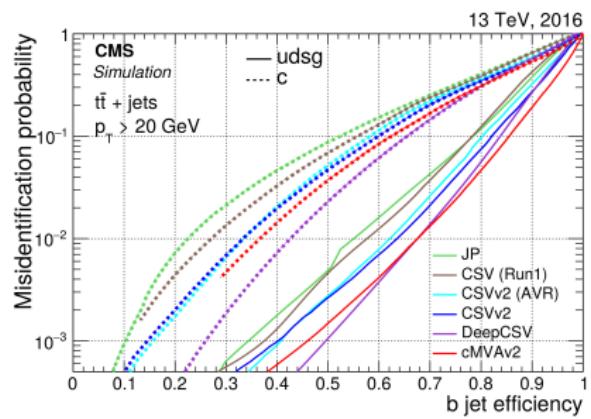
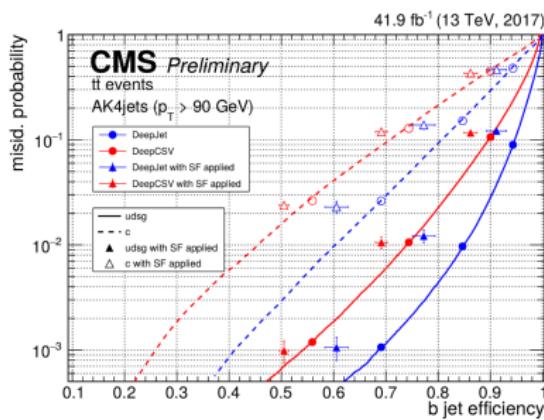
[link]

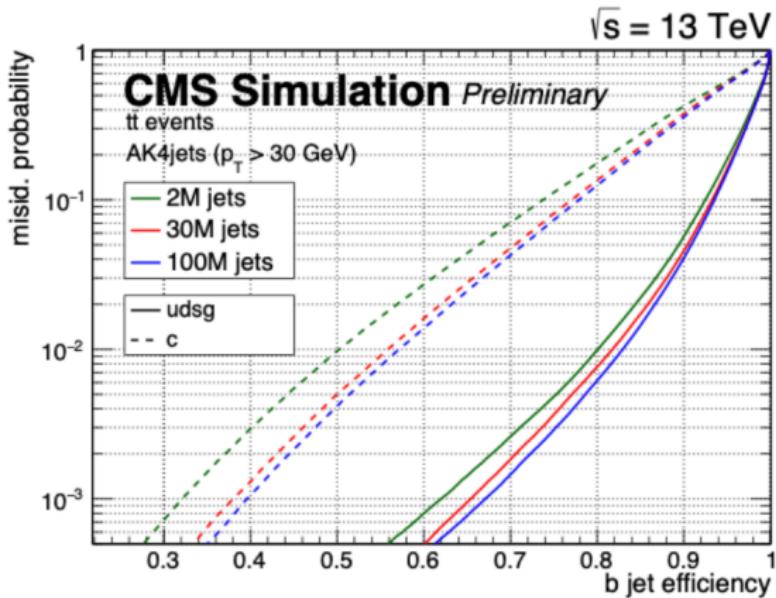
Tagging CMS taggers



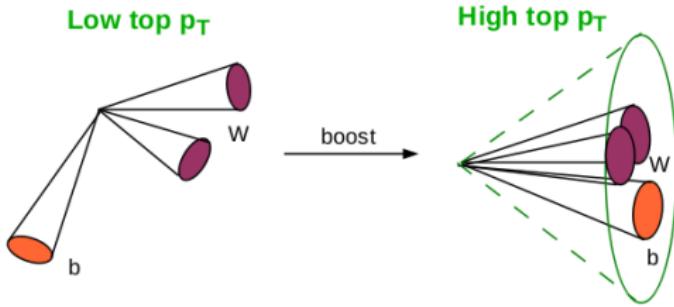
[1], [2]

Tagging ROCs





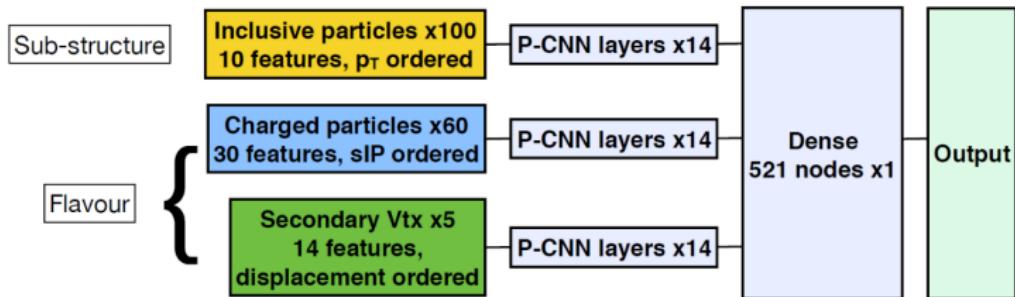
Tagging boosted topology



[link]

Tagging

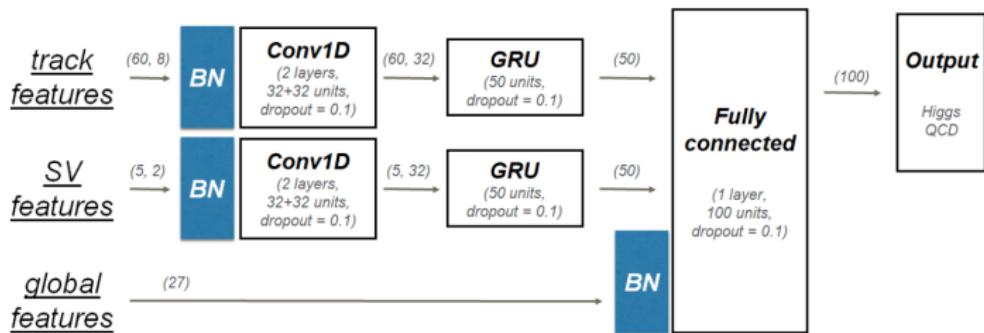
DeepAK8



[link]

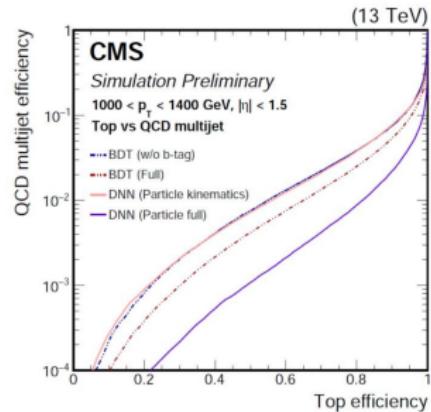
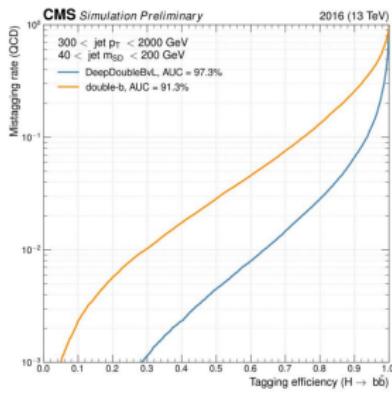
Tagging

DeepDoubleX



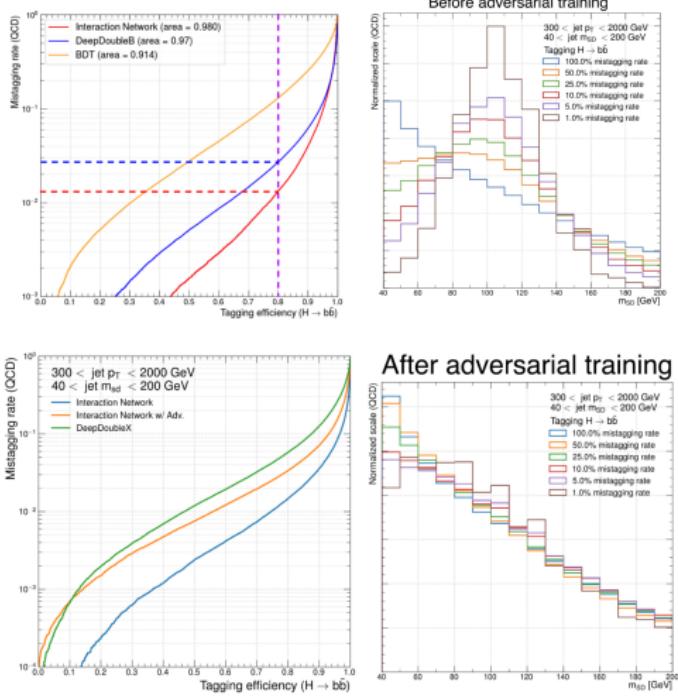
[link]

Tagging boosted comparison



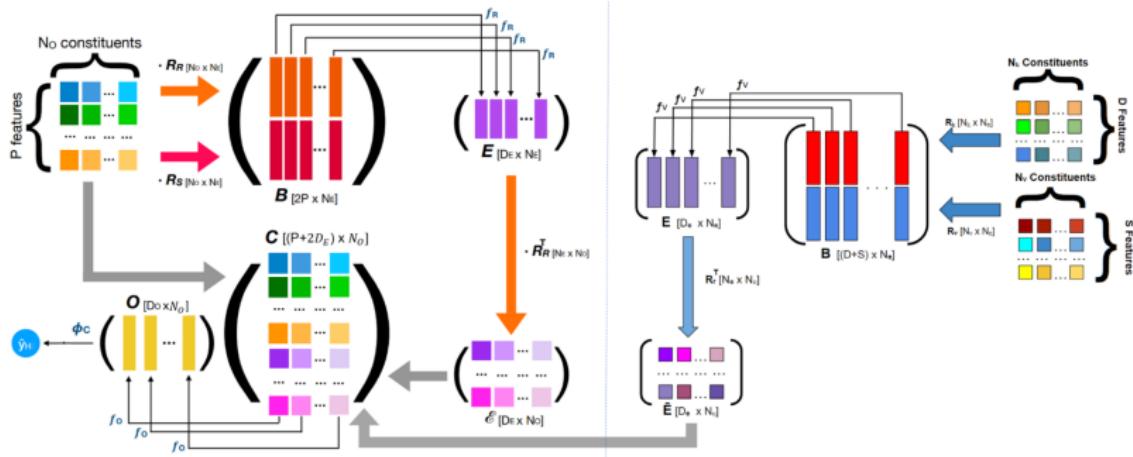
[link]

Tagging Interaction



[link]

Tagging Interaction



Jets

state of the art

Jets

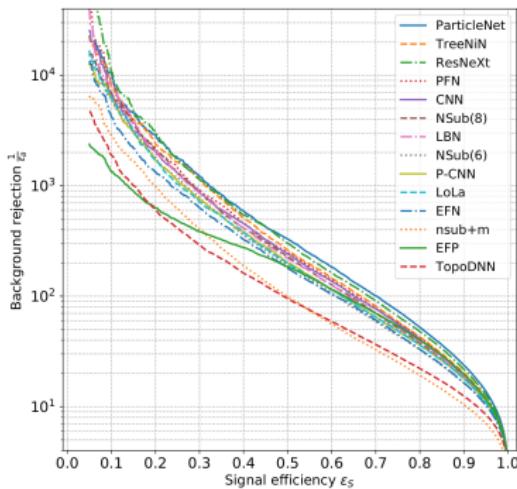
state of the art

(in top tagging)

you should definitely read [this]

State of the Art

top taggers

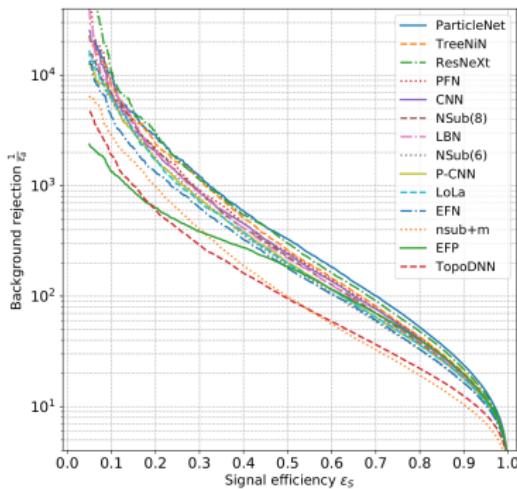


	AUC	Acc	1/ ϵ_B ($\epsilon_S = 0.3$)			#Param
			single	mean	median	
CNN [16]	0.981	0.930	914±14	995±15	975±18	610k
ResNeXt [30]	0.984	0.936	1122±47	1270±28	1286±31	1.46M
TopoDNN [18]	0.972	0.916	295±5	382± 5	378 ± 8	59k
Multi-body N -subjettiness 6 [24]	0.979	0.922	792±18	798±12	808±13	57k
Multi-body N -subjettiness 8 [24]	0.981	0.929	867±15	918±20	926±18	58k
TreeNIN [43]	0.982	0.933	1025±11	1202±23	1188±24	34k
P-CNN	0.980	0.930	732±24	845±13	834±14	348k
ParticleNet [47]	0.985	0.938	1298±46	1412±45	1393±41	498k
LBN [19]	0.981	0.931	836±17	859±67	966±20	705k
LoLa [22]	0.980	0.929	722±17	768±11	765±11	127k
Energy Flow Polynomials [21]	0.980	0.932	384			1k
Energy Flow Network [23]	0.979	0.927	633±31	729±13	726±11	82k
Particle Flow Network [23]	0.982	0.932	891±18	1063±21	1052±29	82k
GoaT	0.985	0.939	1368±140		1549±208	35k

you should definitely read [this]

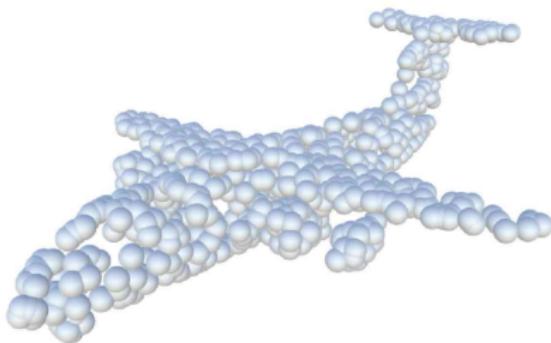
State of the Art

top taggers



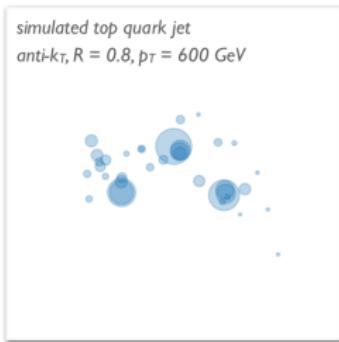
	AUC	Acc	1/ ϵ_B ($\epsilon_S = 0.3$)			#Param
			single	mean	median	
CNN [16]	0.981	0.930	914±14	995±15	975±18	610k
ResNeXt [30]	0.984	0.936	1122±47	1270±28	1286±31	1.46M
TopoDNN [18]	0.972	0.916	295±5	382± 5	378 ± 8	59k
Multi-body N -subjettiness 6 [24]	0.979	0.922	792±18	798±12	808±13	57k
Multi-body N -subjettiness 8 [24]	0.981	0.929	867±15	918±20	926±18	58k
TreeNIN [43]	0.982	0.933	1025±11	1202±23	1188±24	34k
P-CNN	0.980	0.930	732±24	845±13	834±14	348k
ParticleNet [47]	0.985	0.938	1298±46	1412±45	1393±41	498k
LBN [19]	0.981	0.931	836±17	859±67	966±20	705k
LoLa [22]	0.980	0.929	722±17	768±11	765±11	127k
Energy Flow Polynomials [21]	0.980	0.932	384			1k
Energy Flow Network [23]	0.979	0.927	633±31	729±13	726±11	82k
Particle Flow Network [23]	0.982	0.932	891±18	1063±21	1052±29	82k
GoaT	0.985	0.939	1368±140		1549±208	35k

As... POINT CLOUDS?



arXiv:1801.07829

JET AS A ~~POINT~~ CLOUD PARTICLE



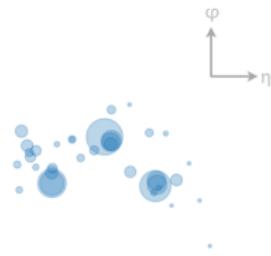
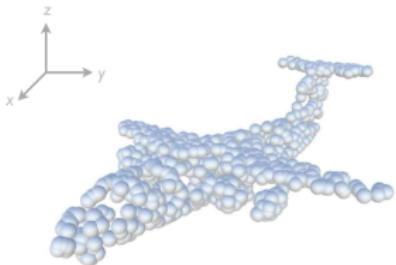
Jet (Particle cloud)

From Wikipedia, the free encyclopedia

A **jet (particle cloud)** is a set of particles in [space](#).

Particle clouds are generally created by clustering a large number of particles measured by [particle detectors](#), e.g.,  and .

POINT CLOUDS VS PARTICLE CLOUDS



- Point cloud

- points are intrinsically *unordered*
 - primary information:
 - 3D coordinates in the xyz space

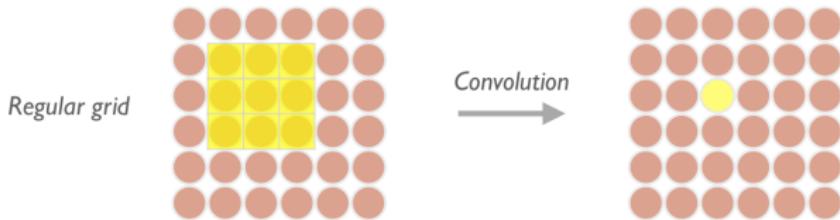
- Particle cloud

- particles are intrinsically *unordered*
 - primary information:
 - 2D coordinates in the $\eta\text{-}\varphi$ space
 - *but also many additional features!*
 - energy/momenta
 - charge/particle ID
 - track quality/impact parameters/etc.

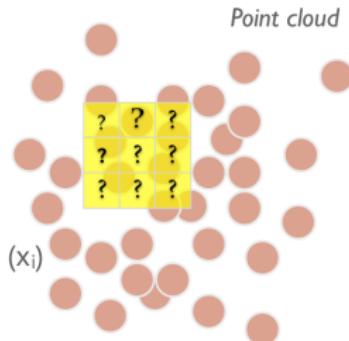
WHY PARTICLE CLOUD?

- Image-based approaches
 - natural idea for calorimeters
 - powerful performance using convolutional neural networks (CNNs)
 - projecting particles into pixels leads to high sparsity and loss of information
 - difficult to include non-additive features (e.g., particle ID, track impact parameters)
- Particle-based approaches
 - preserves full granularity - no loss of information from pixelation
 - straightforward to include any kind of features for each particle
 - constituent particles of a jet are intrinsically unordered – permutation symmetry!
 - however, most particle-based approaches (RNN/TreeNN/1D CNN) failed to exploit this symmetry
- The particle cloud approach
 - also particle-based
 - builds on permutation symmetry from the beginning

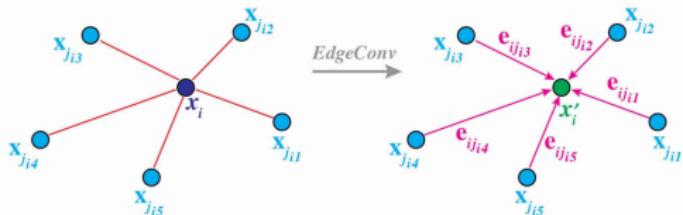
CONVOLUTION ON REGULAR GRIDS



- Conventional convolution only operates on regular grids and cannot be applied on point clouds
 - point clouds are **irregular**
 - how to define a “local” patch to convolve?
 - point clouds are **unordered**
 - conventional convolution operation ($\sum_i K_i x_i$) is not invariant under permutation of the points (x_i)



CONVOLUTION ON POINT CLOUDS

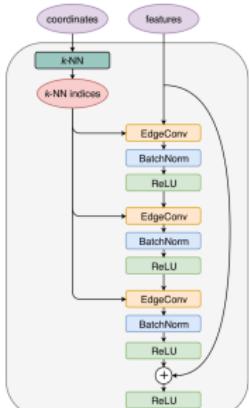


- Convolution on point clouds: **EdgeConv** [arXiv:1801.07829]
 - defining the local patch
 - for each point, a local patch is defined by finding its k-nearest neighbors
 - since points have coordinates, distance can be naturally defined
 - designing a symmetric “convolution” function
 - define “edge feature” for each center-neighbor pair: $e_{ij} = h_\Theta(x_i, x_j)$
 - same h_Θ for all neighbor points, and all center points, for symmetry
 - aggregate the edge features in a symmetric way: $x'_i = \square_j e_{ij}$
 - \square is a permutation-symmetric function, e.g., sum, mean, max

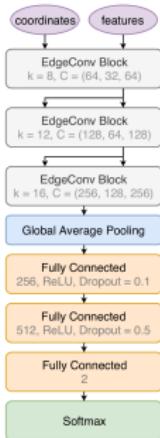
from [link]

State of the Art

ParticleNet



EdgeConv block
(inspired by ResNet)

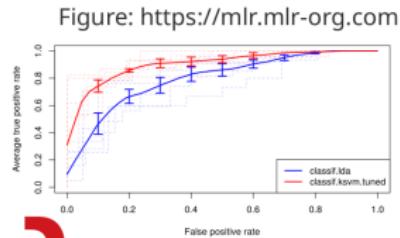


ParticleNet architecture

	AUC	Acc	$1/\epsilon_B$ ($\epsilon_S = 0.3$)	#Param
	single	mean	median	
CNN [16]	0.981	0.930	914±14	995±15
ResNeXt [30]	0.984	0.936	1122±47	1270±28
TopoDNN [18]	0.972	0.916	295±5	378 ± 8
Multi-body N -subjettiness 6 [24]	0.979	0.922	792±18	808±13
Multi-body N -subjettiness 8 [24]	0.981	0.929	867±15	918±20
TreeNIN [43]	0.982	0.933	1025±11	926±18
P-CNN	0.980	0.930	732±24	1188±24
ParticleNet [47]	0.985	0.938	1298±40	1412±45
LBN [19]	0.981	0.931	836±17	859±67
LoLa [22]	0.980	0.929	722±17	768±11
Energy Flow Polynomials [21]	0.980	0.932	384	765±11
Energy Flow Network [23]	0.979	0.927	633±31	726±11
Particle Flow Network [23]	0.982	0.932	891±18	1063±21
Goat	0.985	0.939	1368±140	1549±208

Current ML uncertainties

- 1) Multiple training runs
 - Weight initialisation
- 2) Different training data
 - Training statistical uncertainty
- 3) Different testing data
 - Testing statistical uncertainty

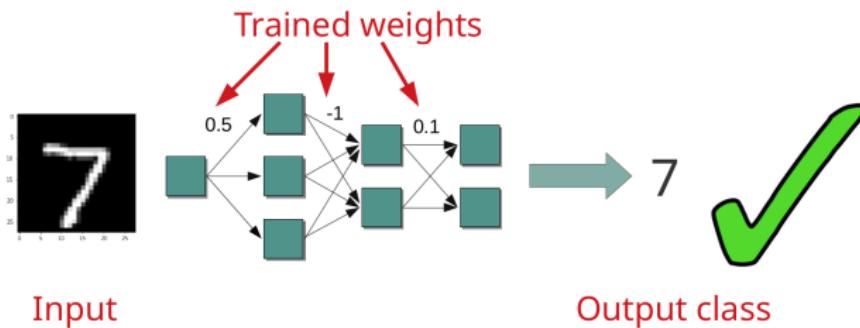


e.g. cross validation

No out-of-distribution or systematic uncertainty

Example: Bayesian motivation

- Example: 1. Train a neural network on MNIST
- 2. Test on MNIST test data
- 3. Success

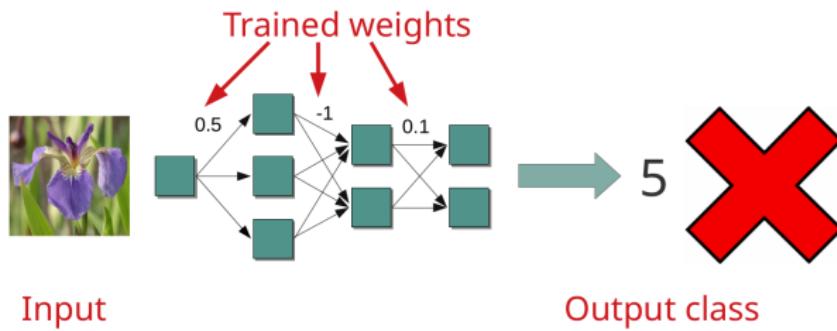


MNIST: LeCun et al., 1998a

4 / 25

Example: Bayesian motivation

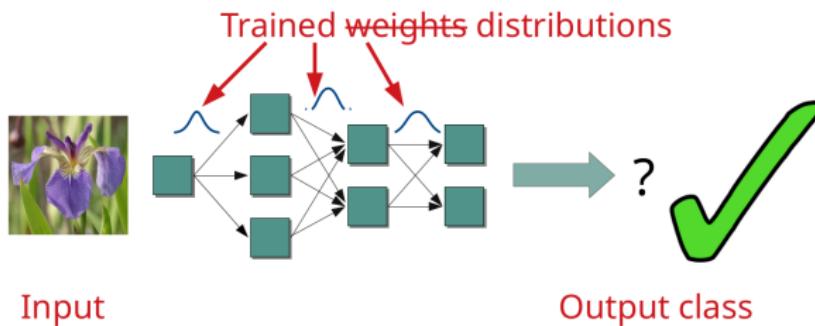
- Example: 1. Train a neural network on MNIST
- 2. Test on ~~MNIST test data~~ new image
- 3. Success Confident, incorrect answer



Iris dataset: doi:10.1111/j.1469-1809.1936.tb02137

Example: Bayesian motivation

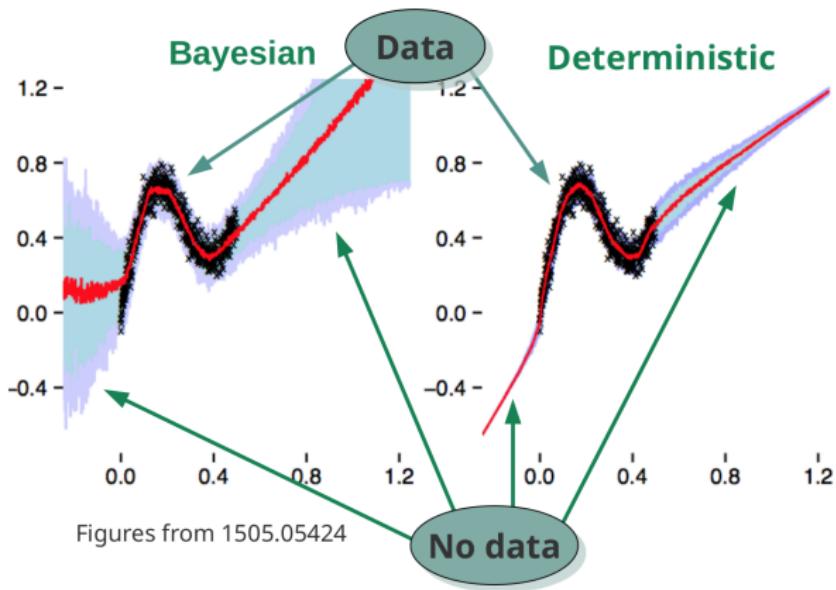
- Example:
1. Train a Bayesian neural network on MNIST
 2. Test on new image
 3. No clear classification
→ Network is unsure



Iris dataset: doi:10.1111/j.1469-1809.1936.tb02137

7 / 25

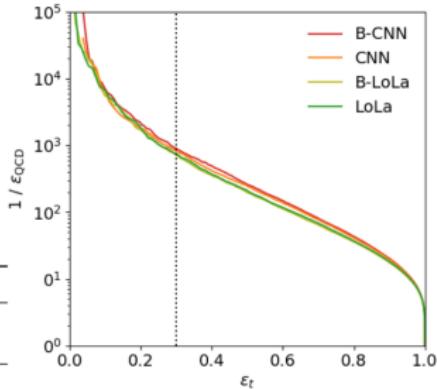
Bayesian networks in practice



Bayesian networks as taggers

Bayesian ~ deterministic
→ State-of-the-art taggers

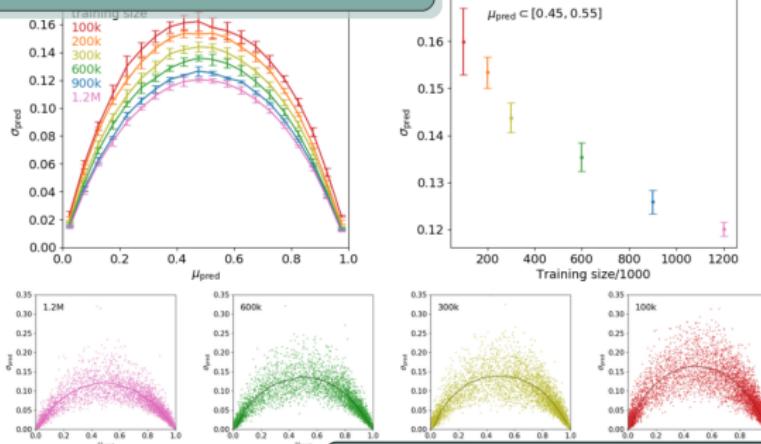
	AUC	$1/\epsilon_{\text{QCD}}$ for $\epsilon_t = 30\%$
CNN	0.982	820
B-CNN	0.982	900
LoLa	$N_{\text{const}} = 40$	0.979
B-LoLa		640
LoLa	$N_{\text{const}} = 100$	0.981
B-LoLa		740
LoLa	$N_{\text{const}} = 200$	0.979
B-LoLa		590



Works with:
- Images (CNN)
- Constituents (LoLa)

Statistical uncertainties

Strong correlation with mean



Training size ↑ uncertainty ↓

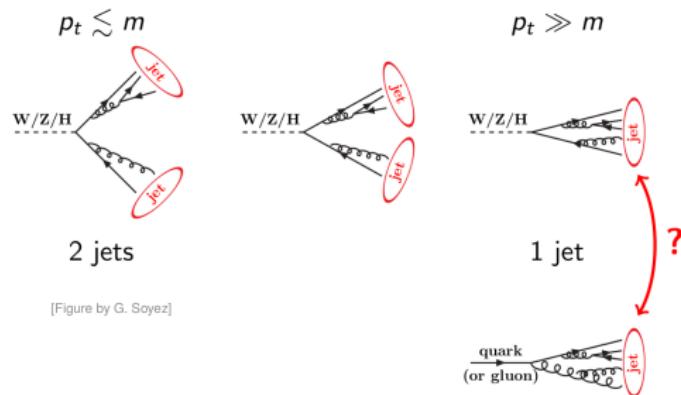
13 / 25

Jets

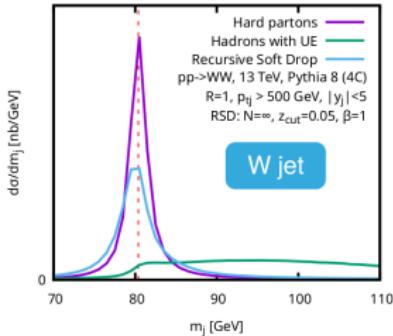
misc.

from [link]

Grooming with RL

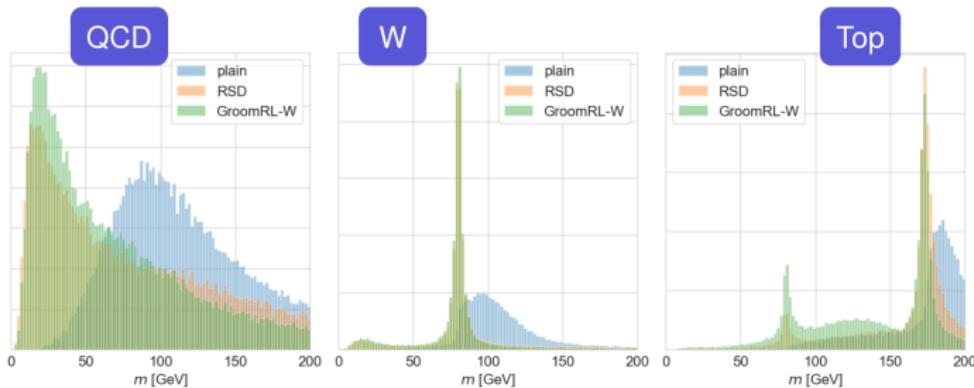
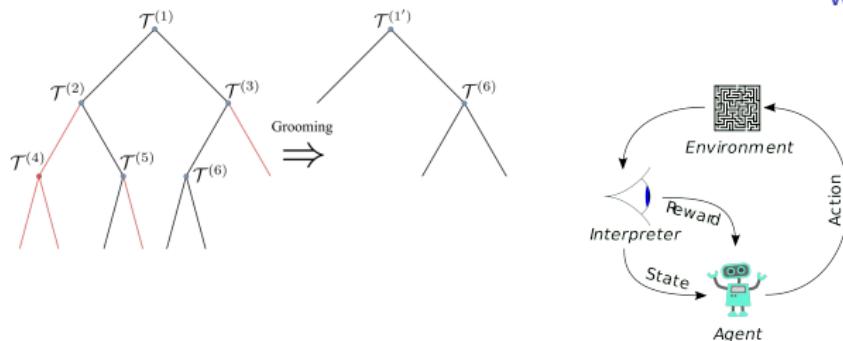


[Figure by G. Soyez]



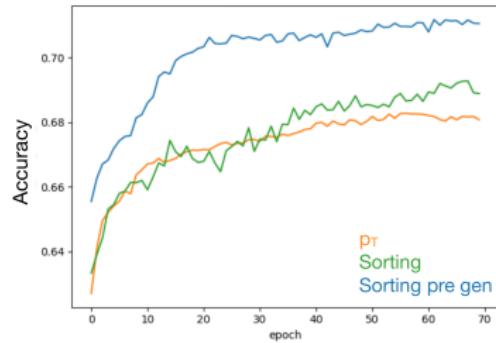
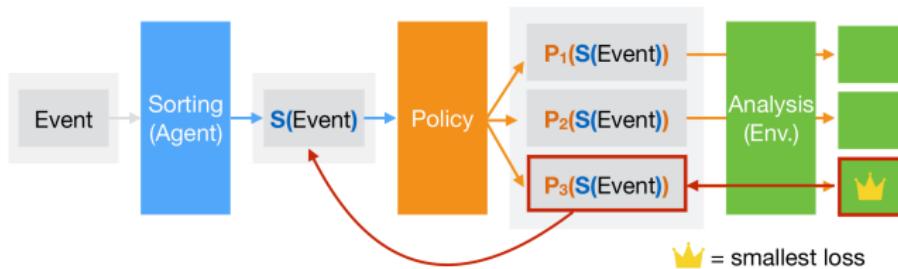
from [link]

Grooming with RL



from [link]

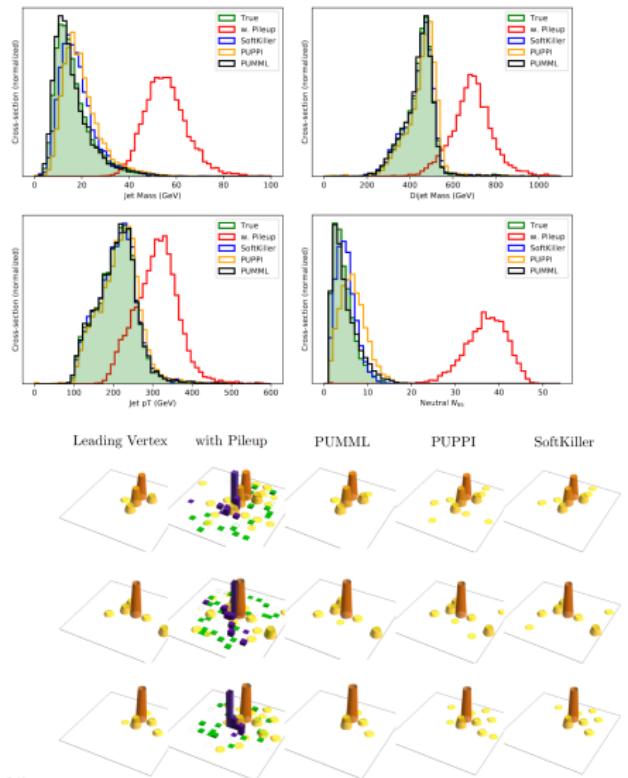
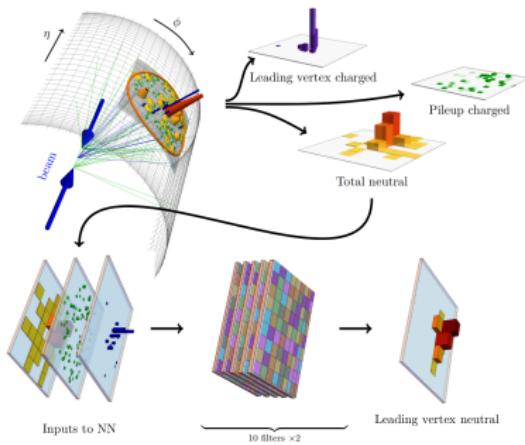
Particle sorting with RL



from [paper]

PU mitigation

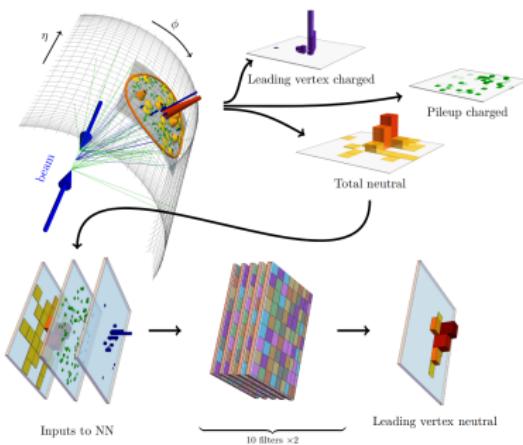
PUMML



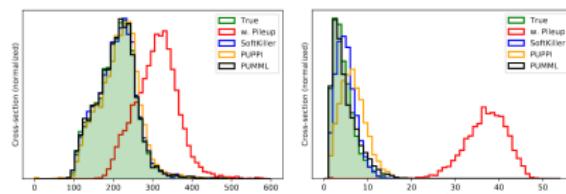
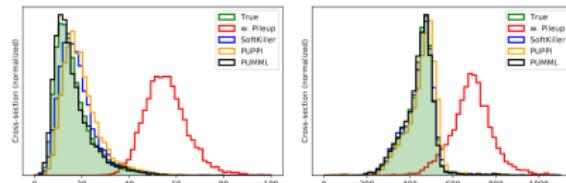
from [paper]

PU mitigation

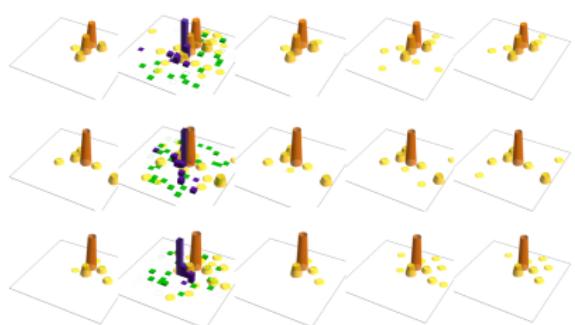
PUMML



which branch of ML?



Leading Vertex with Pileup PUMML PUPPI SoftKiller

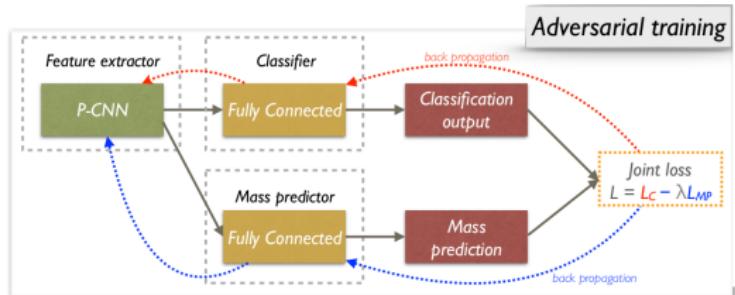
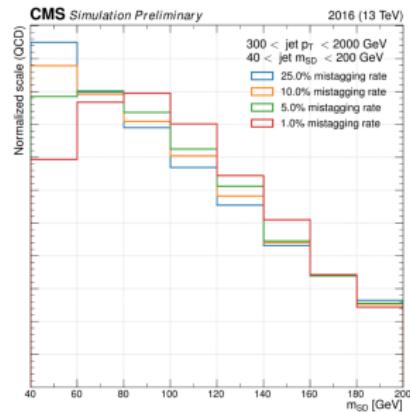


Going adversarial

[link]

Going adversarial

Mass sculpting



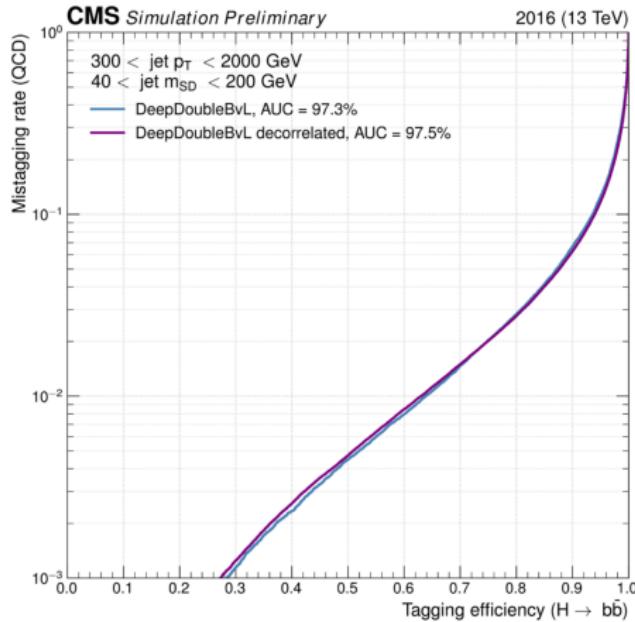
pic from [1]

sculpting in B physics [2]

uBoost [3], and more on mass decorr. [4]

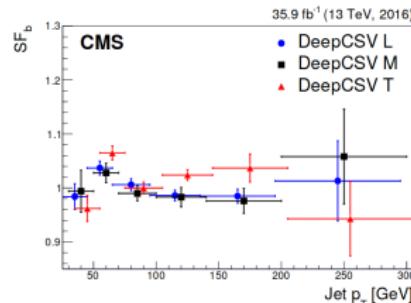
Going adversarial

Mass sculpting



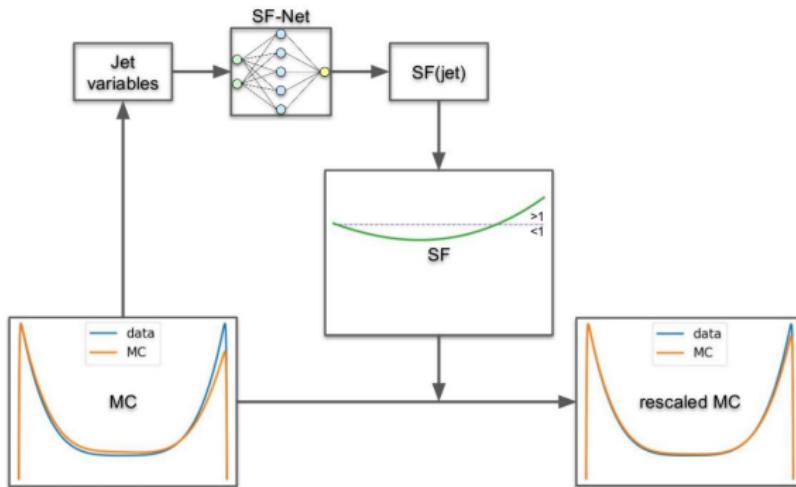
- MC used for training and data differ
 - reweight events to match MC and data
- use Tag&Probe method in bins of p_T, η
 - $SF(p_T, \eta)$

- discrete values
- should have enough statistics in bins
- can't go many variables
- babysitting with bins choice and fitting



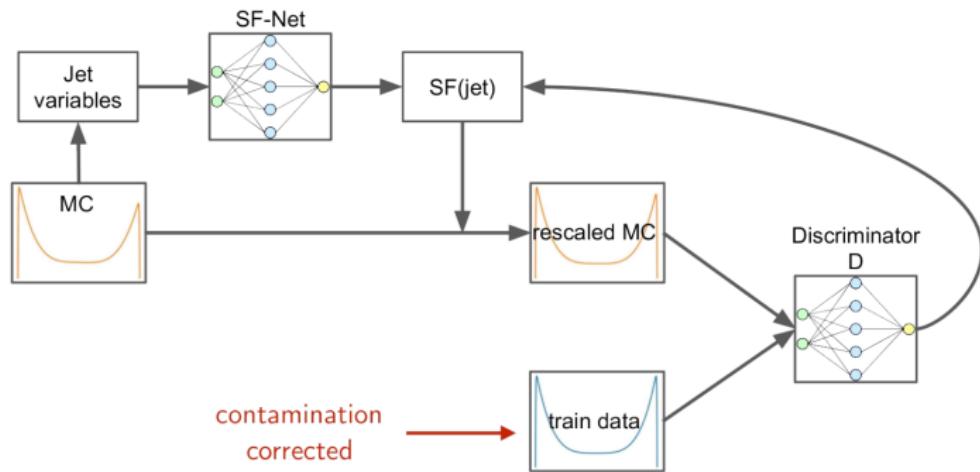
Going adversarial

DeepSF



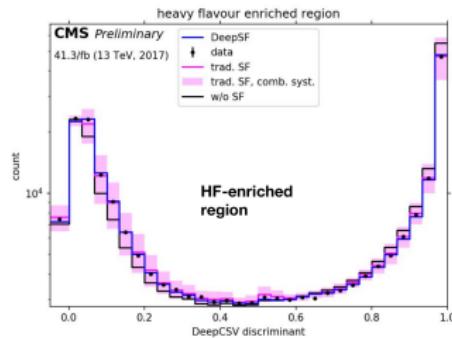
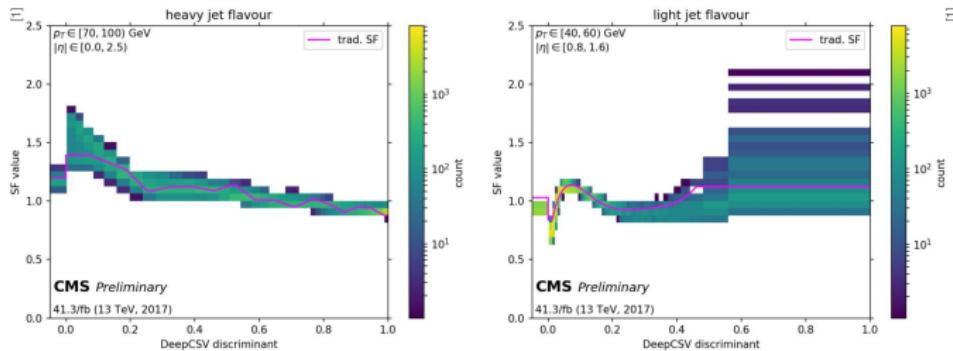
Going adversarial

DeepSF



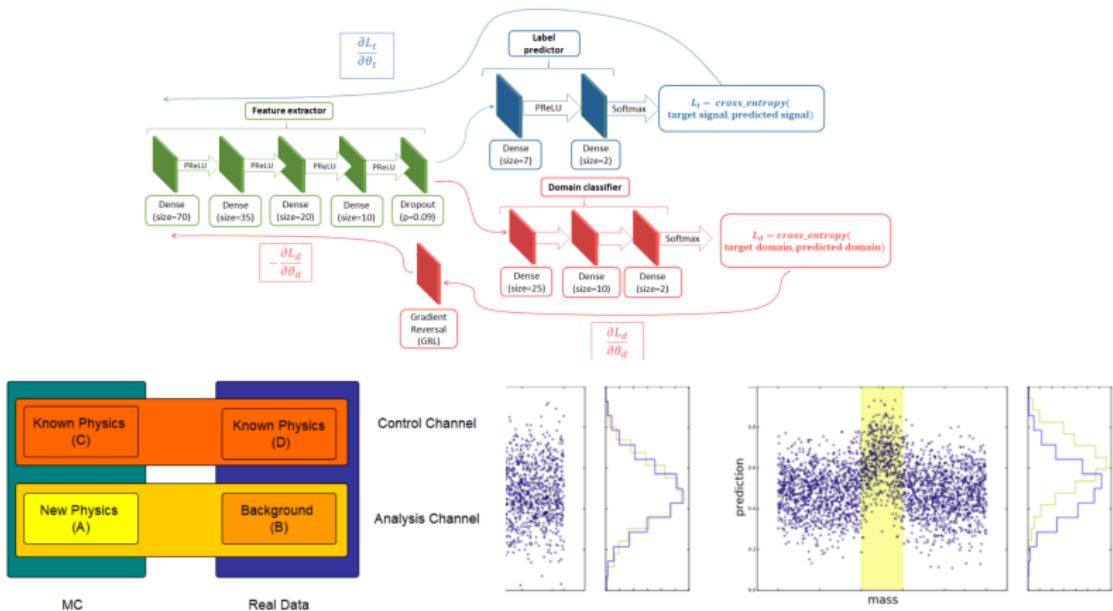
fancy BDT reweighting [link]

Going adversarial DeepSF



more on adversarial methods [1], [2]
Learning to Pivot [3]

Going adversarial Domain adaptation



GAN & VAE

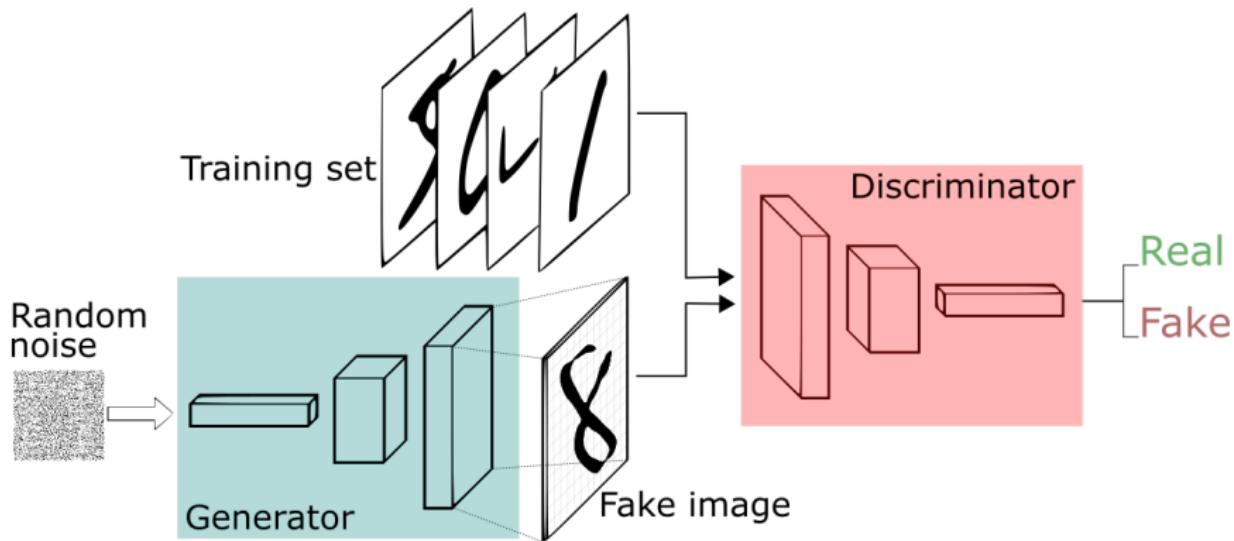
GAN & VAE

Generative Adversarial Network
Variational AutoEncoder

[more here]

GAN

vanilla edition



- GAN - Generative Adversarial Networks
- 3D-GAN - Learning a Probabilistic Latent Space of Object Shapes via 3D Generative-Adversarial Modeling
- acGAN - Face Aging With Conditional Generative Adversarial Networks
- AC-GAN - Conditional Image Synthesis With Auxiliary Classifier GANs
- AdaGAN - AdaGAN: Boosting Generative Models
- AEGAN - Learning Inverse Mapping by Autoencoder based Generative Adversarial Nets
- AffGAN - Amortised MAP Inference for Image Super-resolution
- AL-CGAN - Learning to Generate Images of Outdoor Scenes from Attributes and Semantic Layouts
- ALI - Adversarially Learned Inference
- AM-GAN - Generative Adversarial Nets with Labeled Data by Activation Maximization
- AnoGAN - Unsupervised Anomaly Detection with Generative Adversarial Networks to Guide Marker Discovery
- ArtGAN - Artwork Synthesis with Conditional Categorical GANs
- b-GAN - b-GAN: Unified Framework of Generative Adversarial Networks
- Bayesian GAN - Deep and Hierarchical Implicit Models
- BEGAN - BEGAN: Boundary Equilibrium Generative Adversarial Networks
- BiGAN - Adversarial Feature Learning
- BS-GAN - Boundary-Seeking Generative Adversarial Networks
- CGAN - Conditional Generative Adversarial Nets
- CaloGAN - CaloGAN: Simulating 3D High Energy Particle Showers in Multi-Layer Electromagnetic Calorimeters with Generative Adversarial Networks
- CCGAN - Semi-Supervised Learning with Context-Conditional Generative Adversarial Networks
- CatGAN - Unsupervised and Semi-supervised Learning with Categorical Generative Adversarial Networks
- CoGAN - Coupled Generative Adversarial Networks
- Context-RNN-GAN - Contextual RNN-GANs for Abstract Reasoning Diagram Generation
- C-RNN-GAN - C-RNN-GAN: Continuous recurrent neural networks with adversarial training
- CS-GAN - Improving Neural Machine Translation with Conditional Sequence Generative Adversarial Nets
- CVAE-GAN - CVAE-GAN: Fine-Grained Image Generation through Asymmetric Training
- CycleGAN - Unpaired Image-to-Image Translation using Cycle-Consistent Adversarial Networks
- DTM - Unsupervised Cross-Domain Image Generation
- DCGAN - Unsupervised Representation Learning with Deep Convolutional Generative Adversarial Networks
- DiscoGAN - Learning to Discover Cross-Domain Relations with Generative Adversarial Networks
- DR-GAN - Disentangled Representation Learning GAN for Pose-Invariant Face Recognition
- DualGAN - DualGAN: Unsupervised Dual Learning for Image-to-Image Translation
- EBGAN - Energy-based Generative Adversarial Network
- f-GAN - f-GAN: Training Generative Neural Samplers using Variational Divergence Minimization
- FF-GAN - Towards Large-Pose Face Frontalization in the Wild
- GAWNN - Learning What and Where to Draw
- GeneGAN - GeneGAN: Learning Object Transfiguration and Attribute Subspace from Unpaired Data
- Geometric GAN - Geometric GAN
- GoGAN - Gang of GANs: Generative Adversarial Networks with Maximum Margin Ranking
- GP-GAN - GP-GAN: Towards Realistic High-Resolution Image Blending
- IAN - Neural Photo Editing with Inerspective Adversarial Networks
- iGAN - Generative Visual Manipulation on the Natural Image Manifold
- IcGAN - Invertible Conditional GANs for image editing
- ID-CGAN - Image De-training Using a Conditional Generative Adversarial Network
- Improved GAN - Improved Techniques for Training GANs
- InfoGAN - InfoGAN: Interpretable Representation Learning by Information Maximizing Generative Adversarial Nets
- LAGAN - Learning Particle Physics by Example: Location-Aware Generative Adversarial Networks for Physics Synthesis
- LAPGAN - Deep Generative Image Models using a Laplacian Pyramid of Adversarial Networks

pic from [1]
more fancy [2]

GAN



GAN

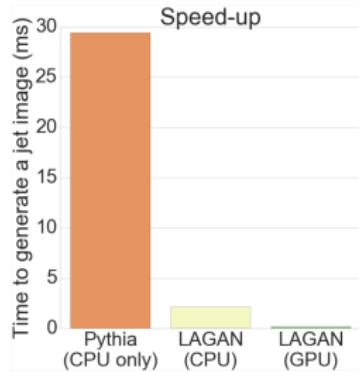
motivation

- **Really** fast simulation

GAN

motivation

- Really fast simulation
in fact $10^2 \div 10^5$ faster than Geant

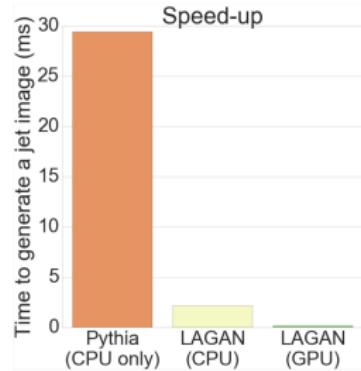


Generation Method	Hardware	Batch Size	milliseconds/shower
GEANT4	CPU	N/A	1772
		1	13.1
		10	5.11
		128	2.19
		1024	2.03
	GPU	1	14.5
		4	3.68
		128	0.021
		512	0.014
		1024	0.012

GAN

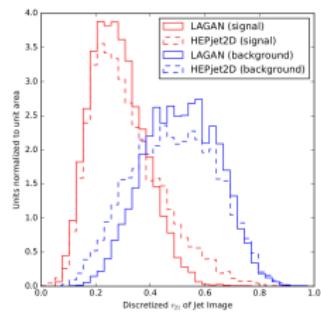
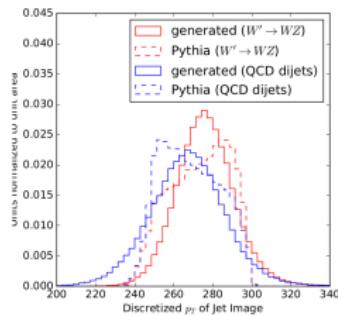
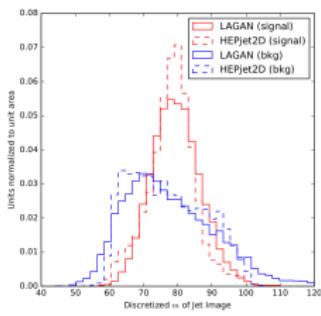
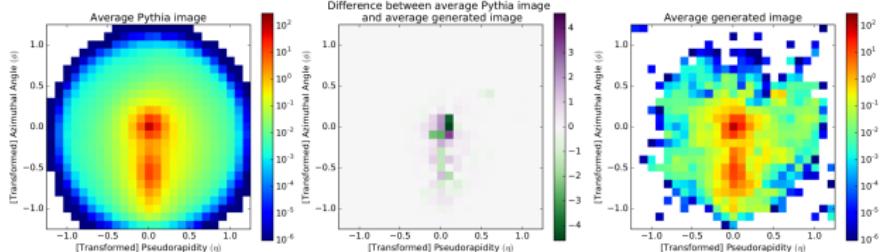
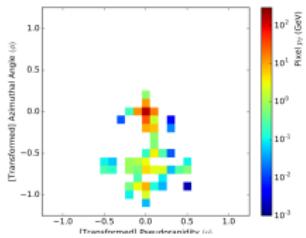
motivation

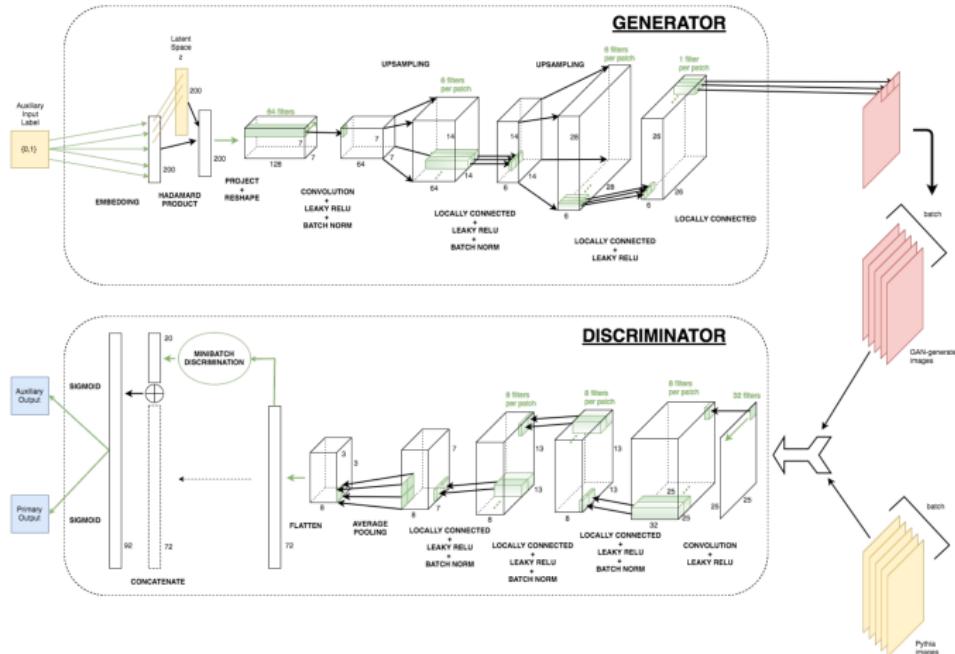
- **Really** fast simulation
in fact $10^2 \div 10^5$ faster than Geant
- Save space
- Rare events



Generation Method	Hardware	Batch Size	milliseconds/shower
GEANT4	CPU	N/A	1772
		1	13.1
		10	5.11
		128	2.19
		1024	2.03
	GPU	1	14.5
		4	3.68
		128	0.021
		512	0.014
		1024	0.012

GAN LAGAN

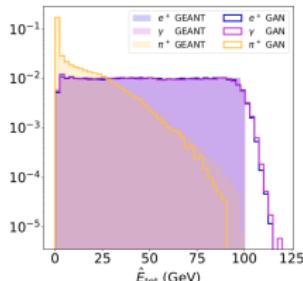
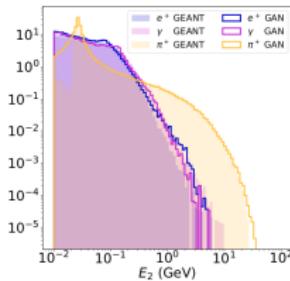
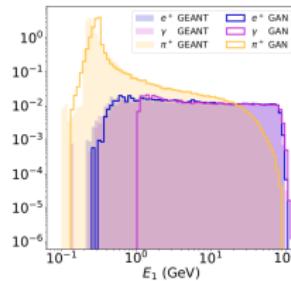
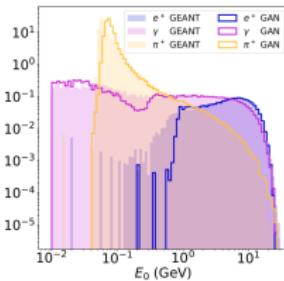
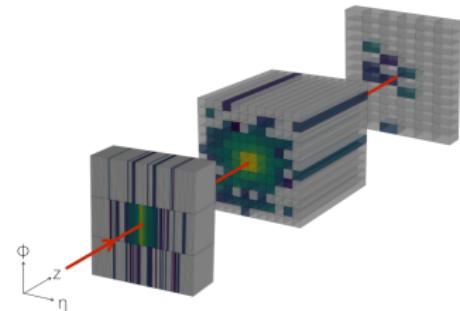
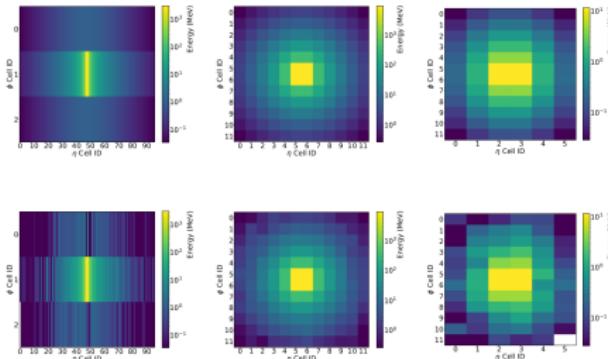




[paper]

GAN

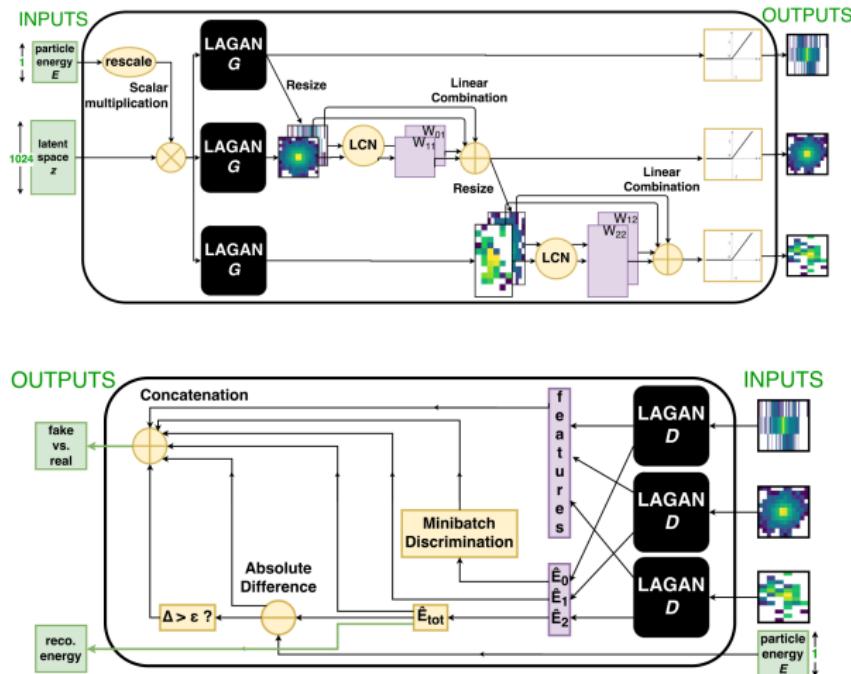
CaloGAN



[paper]

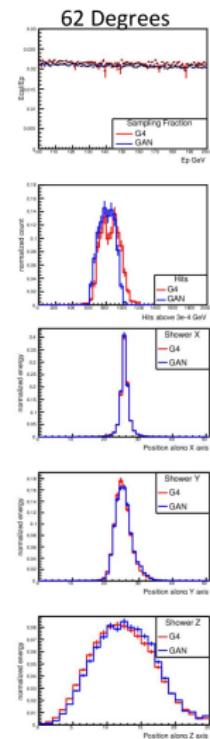
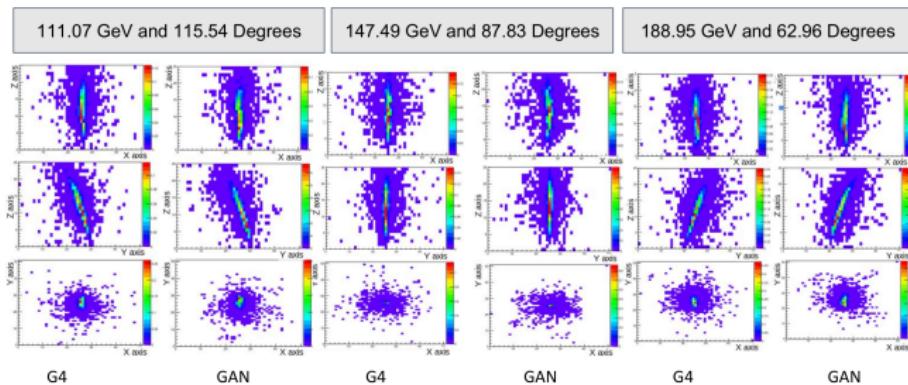
GAN

CaloGAN



[link]

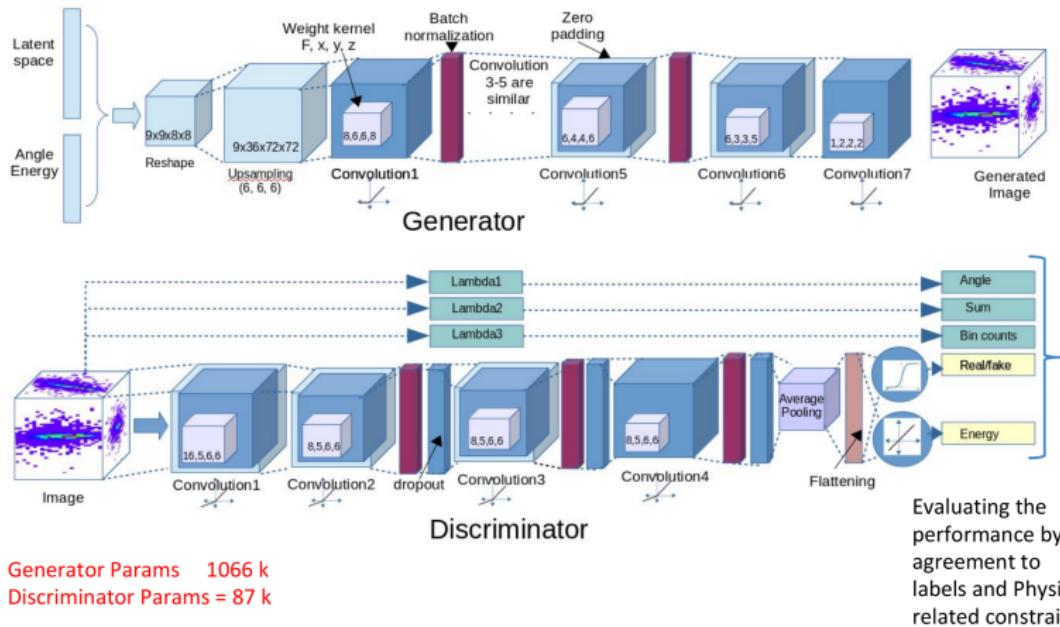
GAN 3DGAN



Time to create an electron shower for Fixed angle		
Method	Machine	Time/Shower (msec)
Full Simulation Fixed Energy (Geant 4)	Intel Xeon Platinum 8180	17000
3DGAN (Fixed Angle) (batch size 128)	Intel Xeon Platinum 8180(TF 1.12)	1
	GTX 1080	0.1
3DGAN (Variable Angle) (batch size 64)	GTX 1080	4.6

[link]

GAN 3DGAN

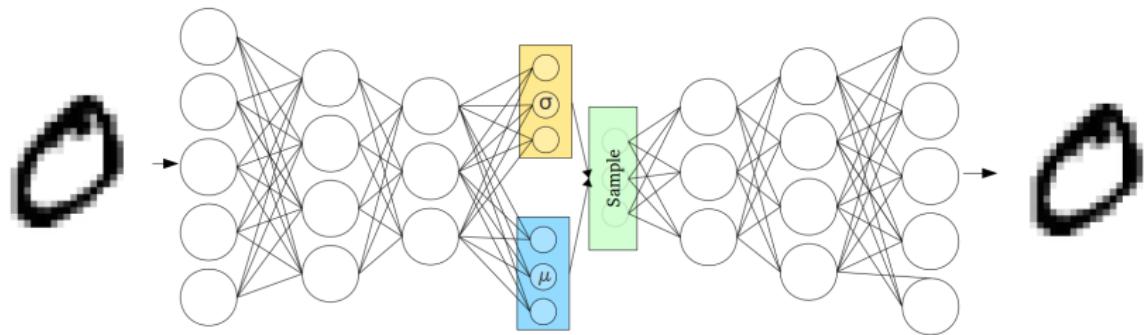


- Generate kinematics
 - DijetGAN (conv) - gen dijet QCD/t \bar{t} [1]
 - Particle-GAN (dense+GRU) - gen min-bias PU? [2]
 - LHCb RICH DLL (dense) [3]
- Generate showers pictures
 - ATLAS ECal GAN (conv) - full ECal geometry accounted for [4]
 - LHCb Ecal GAN (conv) - 2D pics, as usual [5]
 - Belle Pixel bkgr (conv) [6]
 - * GraphRNN + GraphCNN + Mixture Density Network
3D EM showers in SHiP [7]

- Generate kinematics
 - DijetGAN (conv) - gen dijet QCD/t \bar{t} [1]
 - Particle-GAN (dense+GRU) - gen min-bias PU? [2]
 - LHCb RICH DLL (dense) [3]
- Generate showers pictures
 - ATLAS ECal GAN (conv) - full ECal geometry accounted for [4]
 - LHCb Ecal GAN (conv) - 2D pics, as usual [5]
 - Belle Pixel bkgr (conv) [6]
 - * GraphRNN + GraphCNN + Mixture Density Network
3D EM showers in SHiP [7]

what about VAE?

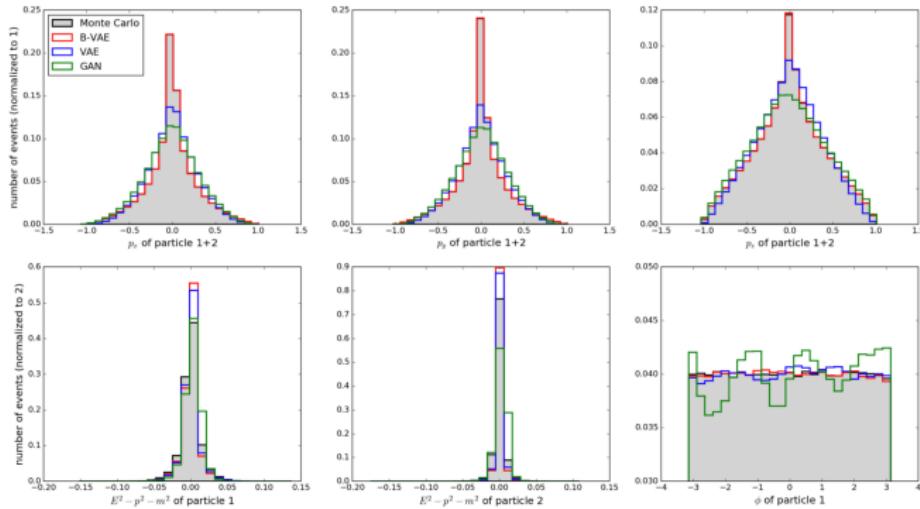
VAE



from [link]

VAE

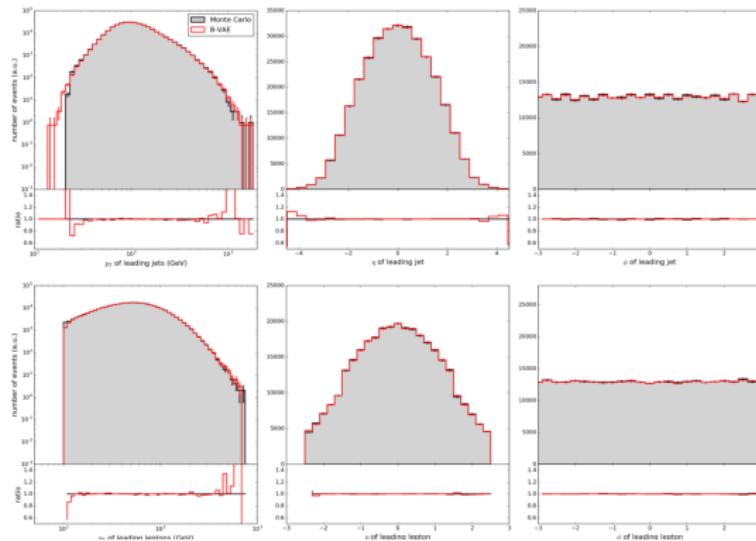
2-body decay



from [link]

VAE

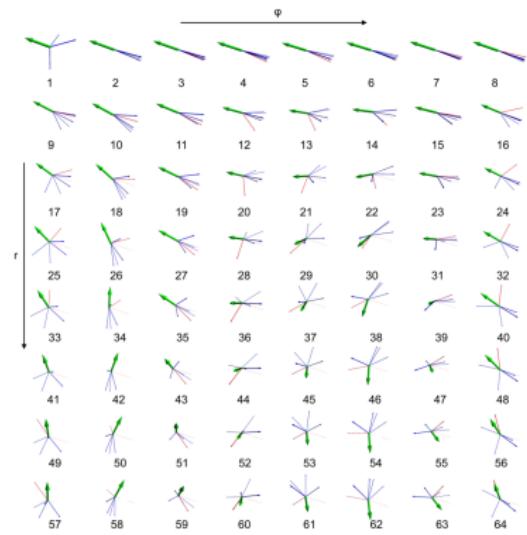
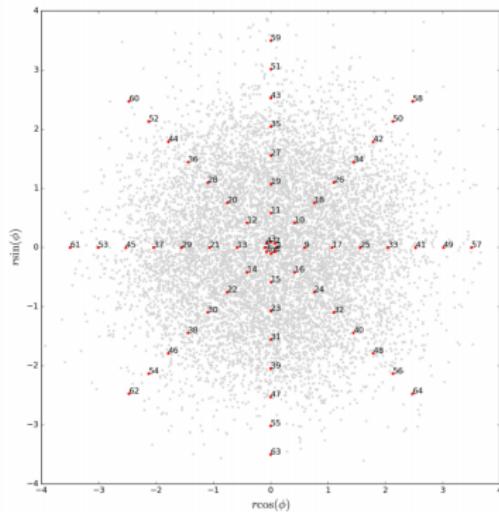
$t\bar{t}$



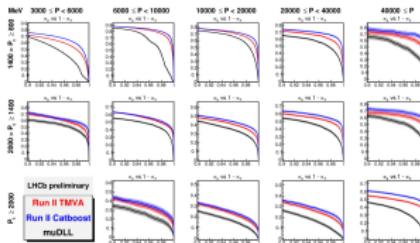
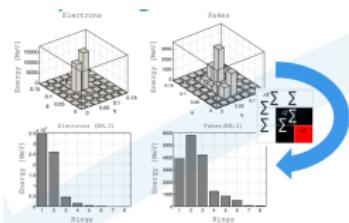
from [link]

VAE

latent space



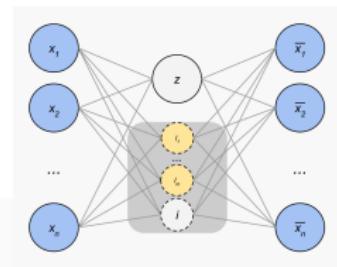
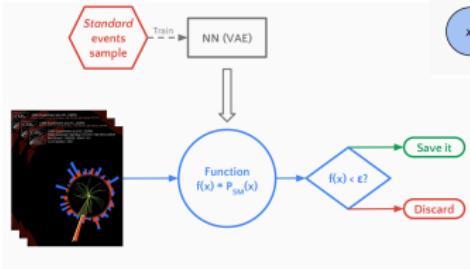
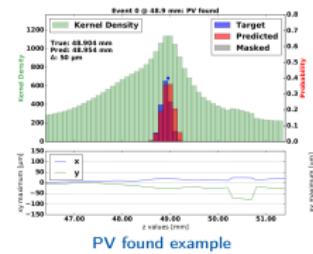
Trigger



Trigger



- **e⁻ HLT in ATLAS** - dense on rings of energy [1]
- **PV finder in LHCb** - CNN on 1D z-KDE + regression [2]
- **hls4ml** - how to deploy DL models on firmware [3]
- **Muon ID in LHCb** - CatBoost [4]
- **Model-independent anomaly detection** - VAE [5]
- **HLT rate monitoring in CMS** - VAE [6]

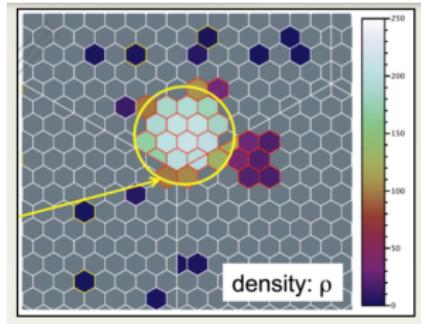
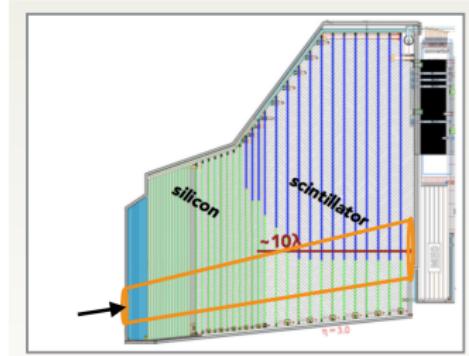


Tracks

from [link]

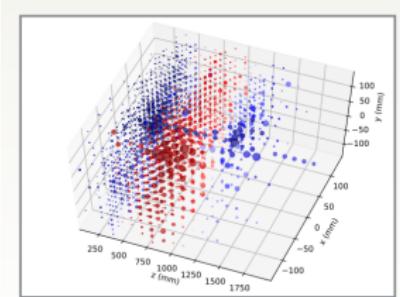
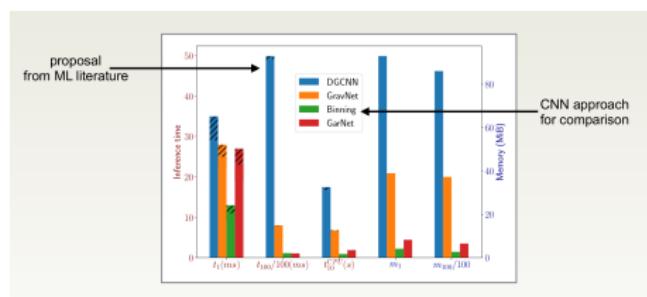
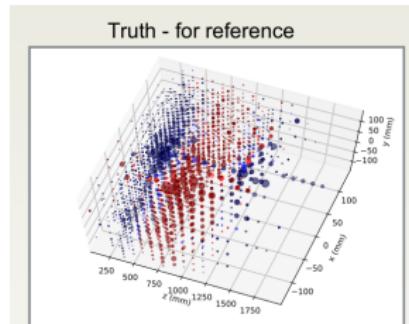
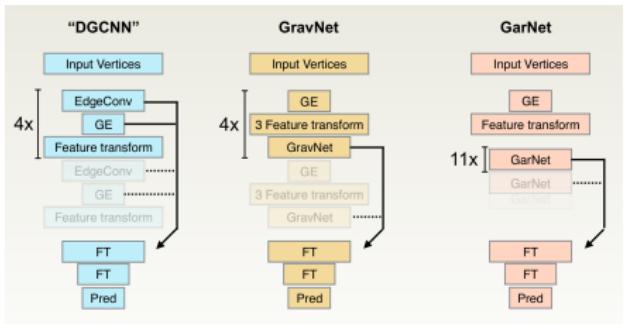
Clustering showers

- ? sparse showers
- ? irregular geometry
- ? CNN
- (good old) point clouds
- modified graph nets - GravNet and GarNet



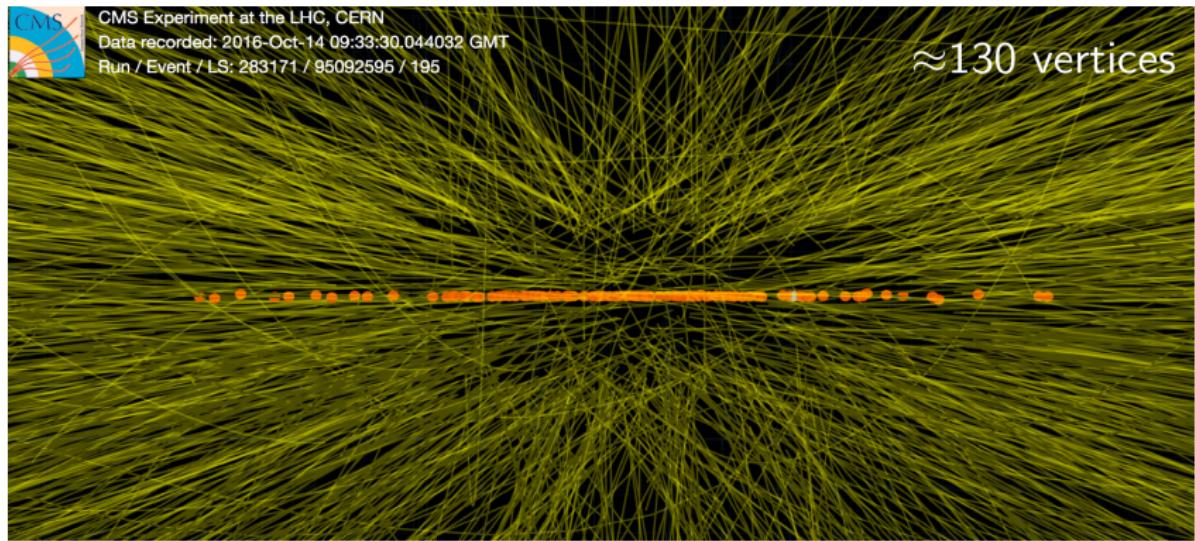
from [link]

Clustering showers



Tracks

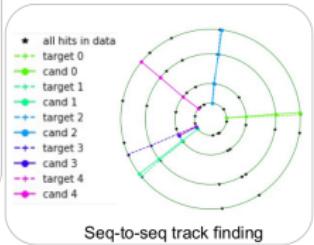
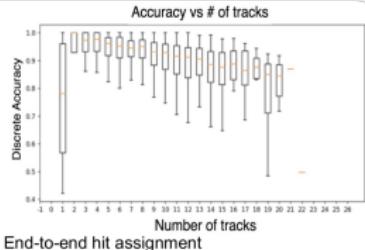
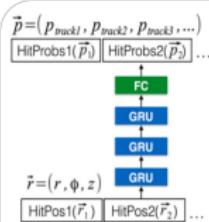
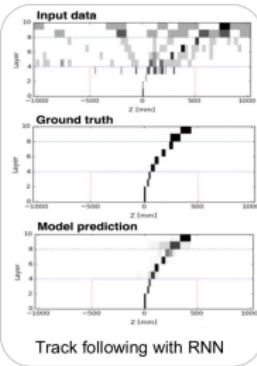
HEP.TrkX



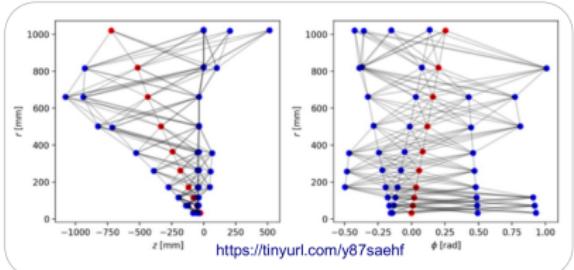
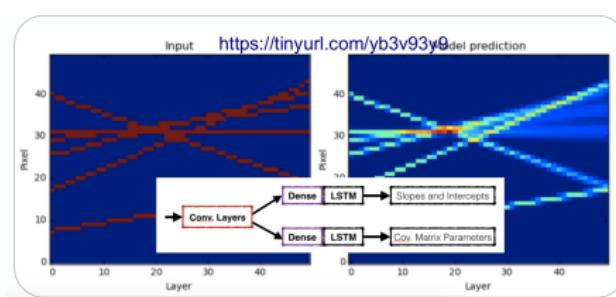
from [link]

Tracks

HEP.TrkX

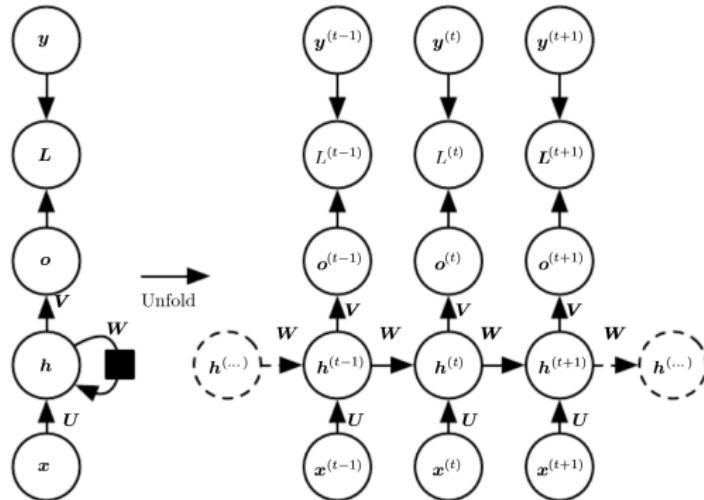


<https://heptrkx.github.io/>



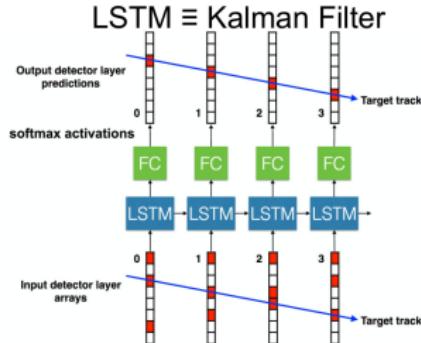
Tracks

HEP.TrkX

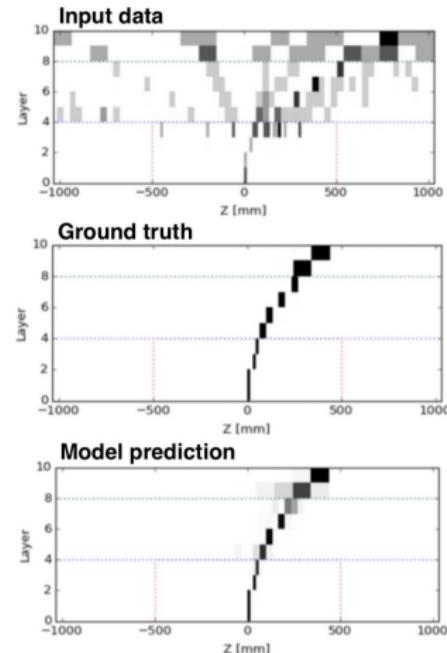


from [link]

Tracks HEP.TrkX



- Search seeded from a known tracklet
- Hit location is discretized to fixed length
- Model predicts the binned position of the hit on the next layer



from [link]

Tracks

trying GNN

- Input Network

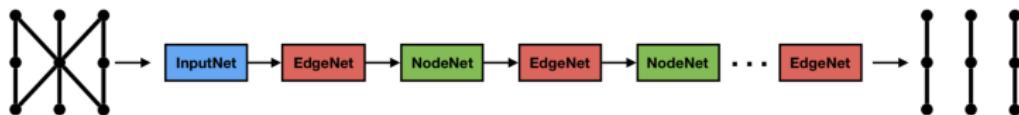
- hit features $(r, \phi, z) \rightarrow$ node latent features
- Dense: $3 \rightarrow N$

- Edge Network

- node representations at both ends \rightarrow edge weight
- Dense: $N + N \rightarrow 1$

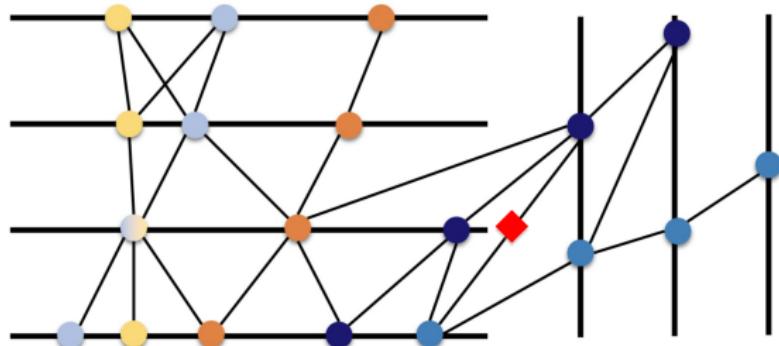
- Node Network

- current node representation, w. sum of incoming edges and w.sum of outgoing edges \rightarrow node latent representation
- Dense: $N + N + N \rightarrow N$



from [link]

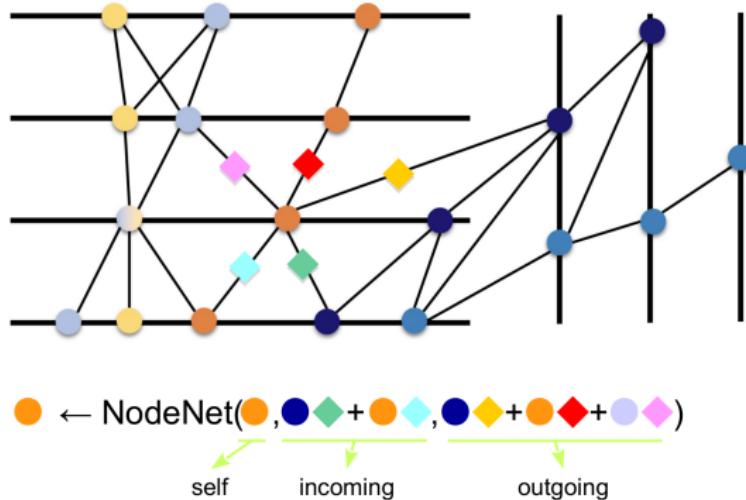
Tracks
trying GNN



◆ ← EdgeNet(●, ●)

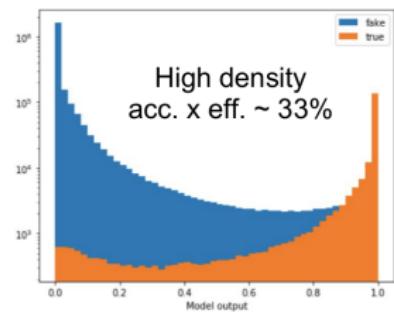
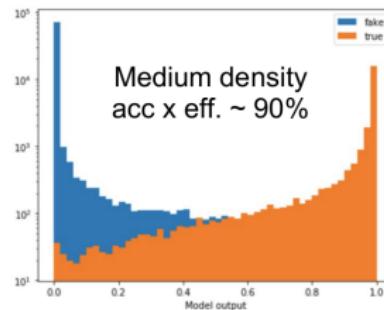
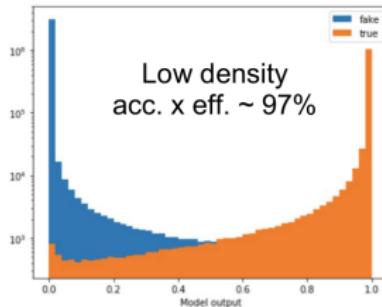
from [link]

Tracks
trying GNN



from [link]

Tracks trying GNN



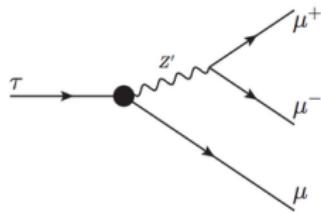
Kaggle

how to compete [1]
nice interview [2]

kaggle

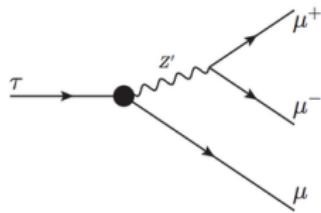
- one of the largest data communities in the world
- ML competitions
- public datasets and kernels (code snippets)
- short-form AI education

- 2014 Higgs Boson challenge by ATLAS [1]
- 2015 $\tau \rightarrow \mu\mu\mu$ by LHCb and Yandex [2]
- 2018 TrackML by CERN [3], [4]



- 2014 Higgs Boson challenge by ATLAS [1]
- 2015 $\tau \rightarrow \mu\mu\mu$ by LHCb and Yandex [2]
- 2018 TrackML by CERN [3], [4]

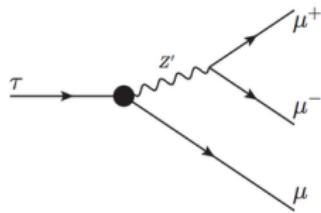
huge impact on ML community in HEP



open data and baselines for practice!

- 2014 Higgs Boson challenge by ATLAS [1]
- 2015 $\tau \rightarrow \mu\mu\mu$ by LHCb and Yandex [2]
- 2018 TrackML by CERN [3], [4]

huge impact on ML community in HEP



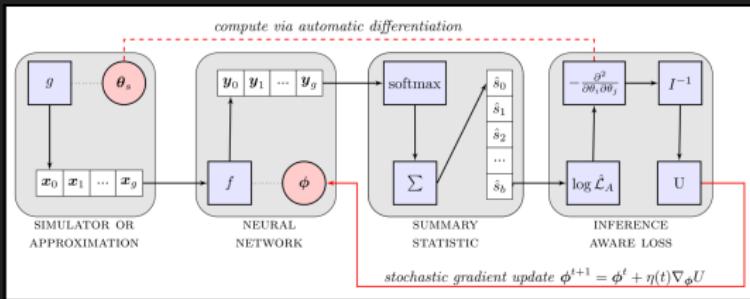
INFERNO

INFERNO

teaser

INFERENCE-AWARE NEURAL OPTIMISATION

An approach to learn non-linear summary statistics by directly minimizing an approximation the expected profiled (or marginalised) interval width accounting for the effect of nuisance parameters



check arxiv.org/abs/1806.04743 for a more detailed description.

INFERNO

if you dare to know

- why **train NN** in the first place
- and why it is **not enough**

INFERNO

if you dare to know

- why **train NN** in the first place
- and why it is **not enough**
- why you really **need MC**

INFERNO

if you dare to know

- why train NN in the first place
- and why it is not enough
- why you really need MC
- what is nuisance parameters
- and why bother with them
- and with test statistic choice

INFERNO

if you dare to know

- why train NN in the first place
- and why it is not enough
- why you really need MC
- what is nuisance parameters
- and why bother with them
- and with test statistic choice
- what is bigger inference perspective
- and why should you care

INFERNO

if you dare to know

- why train NN in the first place [1]
- and why it is not enough
- why you really need MC
- what is nuisance parameters [2]
- and why bother with them
- and with test statistic choice
- what is bigger inference perspective
- and why should you care

Closing remarks

(subjective here)

Closing remarks

future directions

- **interpretability**

- what is NN learning? physics, perhaps?
- do we train it properly?
- can we fully understand it?

- **interpretability**
 - what is NN learning? physics, perhaps?
 - do we train it properly?
 - can we fully understand it?
- **incorporating physics domain knowledge**
 - systematics
 - physics laws and symmetries
 - model's place in HEP pipeline

- **interpretability**
 - what is NN learning? physics, perhaps?
 - do we train it properly?
 - can we fully understand it?
- **incorporating physics domain knowledge**
 - systematics
 - physics laws and symmetries
 - model's place in HEP pipeline
- **close collaboration with ML community**
 - Kaggle
 - industry experience
 - new architectures and concepts

(subjective here)

Closing remarks

about to come

- now

- simple s/b classification in analyses
- particle ID/tagging into production (e.g. CMSSW)
- why don't you use ML in your analysis?

(subjective here)

Closing remarks about to come

- now
 - simple s/b classification in analyses
 - particle ID/tagging into production (e.g. CMSSW)
 - why don't you use ML in your analysis?
- soon
 - ML at trigger level
 - complex taggers (top, boosted) into production
 - centralized simulation with GANs
 - track/object reconstruction
 - architecture competition and evolution (in fact, now)

Closing remarks about to come

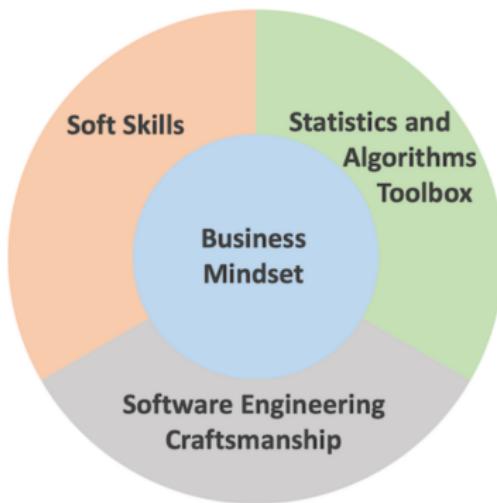
- now
 - simple s/b classification in analyses
 - particle ID/tagging into production (e.g. CMSSW)
 - why don't you use ML in your analysis?
- soon
 - ML at trigger level
 - complex taggers (top, boosted) into production
 - centralized simulation with GANs
 - track/object reconstruction
 - architecture competition and evolution (in fact, now)
- future
 - fully inference/systematics-aware learning
 - HEP specific DL libraries
 - end-to-end training
 - continuous assimilation with ML community
 - ...?

amazing talk
on HEP&ML future
[link]

from [link]

Closing remarks

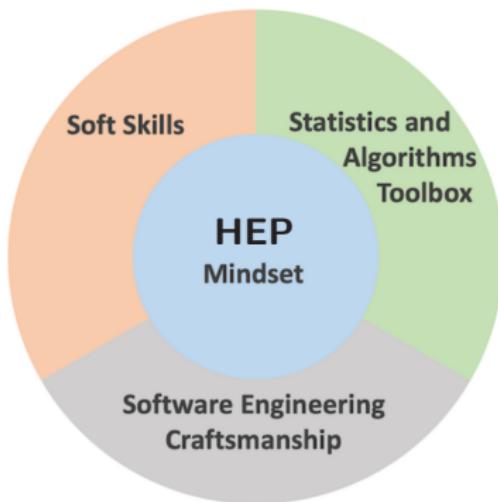
third wave data scientist



from [link]

Closing remarks

third wave data scientist in HEP



Summary

- Jets
 - b-tagging
 - t-tagging
 - grooming&PU
- Going adversarial
 - mass sculpting
 - DeepSF
 - domain adaptation
- Simulating with GAN&VAE
- Trigger
- Tracks
- * Bonus
 - Kaggle
 - INFERNO

@yaourtourto
oleg.filatov@cern.ch

Backup

Артемий Лебедев

Смысл жизни

Смысл жизни в том, чтобы кормить нейросети примерами.

 97K 15:05